

The Founder's Playbook

Building an AI-Native Startup

创始人手册

构建AI原生创业公司

A Guide for the Modern AI-Native Startup Journey

Anthropic / Claude Blog

May 14, 2026

Source: <https://claude.com/blog/the-founders-playbook>

The startup lifecycle,

rebooted for 2026

The 创业公司 lifecycle,

rebooted for 2026

The startup lifecycle, rebooted for 2026

AI is reshaping how startups are built.

Founders who've never written a line of code are shipping production applications today, and the lean 10-person unicorn has gone from scrappy underdog story to deliberate plan of action.

In 2026, AI can write production code, conduct market research, synthesize competitive landscapes, draft investor materials, and automate operational workflows.

By eradicating the once-steep learning curves that even experienced technical founders faced in integrating the tools, platforms, and systems needed to bring their idea to life, AI has above all leveled the playing field around who can launch a startup or build a product.

In 2026, a good idea gets founders further than ever.

Agentic coding compresses what used to take a team of engineers into work a founder can ship themselves.

The traditional startup growth arc assumes that the path from idea to scale is validate → raise → hire → build → raise again → grow → hire more → repeat.

Now, AI has erased the expectation that each new phase in the startup lifecycle requires a bigger team, a different skill set, and a fresh funding round.

This playbook remaps the four core stages of the startup journey (Idea, MVP, Launch, and Scale) according to these new realities.

We examine what each stage looks like when AI is core to your technical and organizational development, what the right tools are for each phase, and how founders using these tools are compressing timelines.

If you're ready to map the shortest path between idea and exit, read on.

2026年创业生命周期的重新定义

AI is reshaping how 创业公司s are built.

Founders who've never written a line of code are shipping 生产环境 applications today, and the lean 10-person unicorn has gone from scrappy underdog story to deliberate plan of action.

In 2026, AI can write 生产环境 code, conduct market 研究/调研, synthesize competitive landscapes, draft investor materials, and automate operational workflows.

By eradicating the once-steep learning curves that even experienced technical 创始人s faced in integrating the tools, platforms, and systems needed to bring their idea to life, AI has above all leveled the playing field around who can 发布/上线 a 创业公司 or build a product.

In 2026, a good idea gets 创始人s further than ever.

代理编程 compresses what used to take a team of engineers into work a 创始人 can ship themselves.

The traditional 创业公司 growth arc assumes that the path from idea to 规模化 is validate → raise → hire → build → raise again → grow → hire more → repeat.

Now, AI has erased the expectation that each new phase in the 创业公司 lifecycle requires a bigger team, a different skill set, and a fresh funding round.

This playbook remaps the four core stages of the 创业公司 journey (Idea, 最小可行产品 (MVP), Launch, and Scale) according to these new realities.

We examine what each stage looks like when AI is core to your technical and organizational development, what the right tools are for each phase, and how 创始人s using these tools are compressing timelines.

If you're ready to map the shortest path between idea and exit, read on.

What it means to be a

founder is changing

What it means to be a

创始人 is changing

What it means to be a founder is changing

Founders used to be defined by what they could do: technical founders wrote code, non-technical founders ran business ops and closed deals.

But the models, systems, and AI agents available to founders in 2026 have dissolved the wall between "people who can build" and "people with ideas worth building."

"AI-native startups are fundamentally transforming what it means to be a founder."

Now someone with no engineering background can build production software that brings their idea to life, while a technically adept founder with little business knowledge can easily produce a go-to-market strategy, a financial model, and a highly polished pitch deck.

Historically, founders spent the bulk of their time in execution mode: writing code, managing people, handling day-to-day operational work.

In an AI-native startup, the founder role becomes much less individual contributor and much more orchestrator of agents—specialized AI assistants that can read files, run commands, execute code, and even browse the web.

The founder's attention shifts up the stack toward the higher-order work: generating ideas and directing the systems (AI agents, tools, and whatever small team exists) that carry those ideas out.

The most revolutionary result of AI as central infrastructure, though, is to unblock non-technical founders with subject matter expertise.

When the founding pool expands beyond people with engineering backgrounds, you get startups built by people with radically different lived experiences, solving real problems that the traditional tech-founder pipeline never prioritized (or perhaps even noticed).

成为创业者的意义正在改变

Founders used to be defined by what they could do: technical 创始人s wrote code, non-technical 创始人s ran business ops and closed deals.

But the models, systems, and AI agents available to 创始人s in 2026 have dissolved the wall between "people who can build" and "people with ideas worth building."

"AI原生 创业公司s are fundamentally transforming what it means to be a 创始人."

Now someone with no engineering background can build 生产环境 software that brings their idea to life, while a technically adept 创始人 with little business knowledge can easily produce a go-to-market strategy, a financial model, and a highly polished 路演PPT.

Historically, 创始人s spent the bulk of their time in 执行 mode: writing code, managing people, handling day-to-day operational work.

In an AI原生 创业公司, the 创始人 role becomes much less individual contributor and much more orchestrator of agents—specialized AI assistants that can read files, run commands, execute code, and even browse the web.

The 创始人's attention shifts up the stack toward the higher-order work: generating ideas and directing the systems (AI agents, tools, and whatever small team exists) that carry those ideas out.

The most revolutionary result of AI as central 基础设施, though, is to unblock non-technical 创始人s with subject matter expertise.

When the founding pool expands beyond people with engineering backgrounds, you get 创业公司s built by people with radically different lived experiences, solving real problems that the traditional tech-创始人 流水线/管道 never prioritized (or perhaps even noticed).

AI tool capabilities for lean startups

The traditional startup model assumed you needed to hire engineers to build, salespeople to sell, and ops people to run the business.

Headcount was treated as a sign of organizational momentum and product maturity.

Early-stage startups in 2026 are radically different.

They're extremely lean by design, often just the founder alone or a team with a few others.

By centering both technical and organizational development on AI as infrastructure, they can reach product validation, early revenue, or even profitability before scaling the team.

There are three areas in particular where AI helps a startup function like a much larger org: research, agentic coding, and automating workflows for key business operations.

轻量级创业公司的AI工具能力

The traditional 创业公司 model assumed you needed to hire engineers to build, salespeople to sell, and ops people to run the business.

Headcount was treated as a sign of organizational momentum and product maturity.

Early-stage 创业公司s in 2026 are radically different.

They're extremely lean by design, often just the 创始人 alone or a team with a few others.

By centering both technical and organizational development on AI as 基础设施, they can reach product 验证, early 营收, or even 盈利 before scaling the team.

There are three areas in particular where AI helps a 创业公司 function like a much larger org: 研究/调研, 代理编程, and automating workflows for key business operations.

Conversational intelligence and research

Think: on-call expert for every domain Consider everything a founder needs to know in the first year that they almost certainly don't know going in: how do I set up payroll?

How do I plan product development sprints? How do I draft a tight investor memo?

Early-stage startup questions like these all used to have the same answer, which was Find someone who knows.

For a bootstrapped or pre-seed founder, this could consume time spent knowledge-gathering instead of building, or possibly requiring burning a chunk of early capital on a consultant.

Now, they have AI as an on-call expert across every conceivable domain.

- Deep research: competitive analysis, market sizing, financial modeling
- Document drafting: pitch decks, case studies, investor memos, PRDs
- Strategic thinking partner: devil's advocate analysis, pre-mortems, scenario

planning, roadmap optimization

对话式智能与研究

Think: on-call expert for every domain Consider everything a 创始人 needs to know in the first year that they almost certainly don't know going in: how do I set up payroll?

How do I plan product development sprints?

How do I draft a tight investor memo?

Early-stage 创业公司 questions like these all used to have the same answer, which was Find someone who knows.

For a bootstrapped or pre-seed 创始人, this could consume time spent knowledge-gathering instead of building, or possibly requiring burning a chunk of early capital on a consultant.

Now, they have AI as an on-call expert across every conceivable domain.

Deep 研究/调研: 竞争分析, market sizing, financial modeling

Document drafting: 路演PPTs, case studies, investor memos, PRDs

Strategic thinking partner: devil's advocate analysis, pre-mortems, scenario

planning, 路线图 optimization

Agentic coding

Think: the engineer who's always available, never blocked
Building software used to require a technical co-founder, a contract dev shop, or a long enough runway to hire an engineering team before you'd written a line of production code.

Agentic coding tools now allow every aspiring founder to describe what they want to build in plain language and direct AI to generate, test, debug, and refactor a production-grade codebase at the speed and scale of a full engineering team.

The timeline from “I have an idea” to “I have a product” has compressed.

And the founder's role now centers on what to build and why, while AI handles the actual construction of real infrastructure that's ready for real users.

代理编程

Think: the engineer who's always available, never blocked
Building software used to require a technical co-创始人, a contract dev shop, or a long enough runway to hire an engineering team before you'd written a line of 生产环境 code.

代理编程 tools now allow every aspiring 创始人 to describe what they want to build in plain language and direct AI to generate, test, 调试, and refactor a 生产环境-grade codebase at the speed and 规模化 of a full engineering team.

The timeline from “I have an idea” to “I have a product” has compressed.

And the 创始人's role now centers on what to build and why, while AI handles the actual construction of real 基础设施 that's ready for real users.

Workflow automation

Think: on-demand, automated ops team Even when a founder can research like a consultant and build like an engineering team, there's still a whole category of work beyond strategic planning or product development that still has to get done.

Scheduling, updating the CRM, pulling weekly reports, keeping documentation current, publishing content, tracking compliance requirements, managing the connective tissue between the tools and systems the company runs on all have to happen, too.

In a lean startup, this load falls mainly on the founder—and it's a significant tax on the time and attention that should be going toward higher-order decisions.

Workflow automation with AI tools offloads that tax.

Recurring operational tasks can be configured to happen automatically so that the CRM updates when a deal moves, a weekly report compiles itself, and product documentation gets updated in sync with product changes.

And, crucially, Claude Cowork integrates with the interconnected systems a startup runs on—your project management tool, your communication stack, your data sources—without needing someone to build and maintain those integrations.

In Day Zero startups, that someone is almost always the founder.

Timing and orchestration are everything Founders that effectively harnesses AI's research, automation, and agentic coding capabilities can build a startup that operates with far more leverage than its headcount suggests.

They also get to dedicate the majority of their time and bandwidth to the work that actually matters.

This work doesn't happen on autopilot; the founder orchestrating these AI tools needs to know how (and when) to apply them.

The rest of this playbook is dedicated to exploring the goals and challenges founders will encounter as they follow the AI-native startup path, and how to effectively apply AI tools at each stage of the journey.

工作流自动化

Think: on-demand, automated ops team Even when a 创始人 can 研究/调研 like a consultant and build like an engineering team, there's still a whole category of work beyond strategic planning or product development that still has to get done.

Scheduling, updating the CRM, pulling weekly reports, keeping documentation current, publishing content, tracking compliance requirements, managing the connective tissue between the tools and systems the company runs on all have to happen, too.

In a lean 创业公司, this load falls mainly on the 创始人—and it's a significant tax on the time and attention that should be going toward higher-order decisions.

工作流自动化 with AI tools offloads that tax.

Recurring operational tasks can be configured to happen automatically so that the CRM updates when a deal moves, a weekly report compiles itself, and product documentation gets updated in sync with product changes.

And, crucially, Claude Cowork integrates with the interconnected systems a 创业公司 runs on—your project management tool, your communication stack, your data sources—without needing someone to build and maintain those 集成s.

In Day Zero 创业公司s, that someone is almost always the 创始人.

时间与编排 are everything Founders that effectively harnesses AI's 研究/调研, automation, and 代理编程 capabilities can build a 创业公司 that operates with far more leverage than its 团队规模 suggests.

They also get to dedicate the majority of their time and bandwidth to the work that actually matters.

This work doesn't happen on autopilot; the 创始人 orchestrating these AI tools needs to know how (and when) to apply them.

The rest of this playbook is dedicated to exploring the goals and challenges 创始人s will encounter as they follow the AI原生 创业公司 path, and how to effectively apply AI tools at each stage of the journey.

Idea Stage

Every startup founder starts from the same place: a problem they can't stop thinking about.

This is the startup phase where idea meets reality: startup success in 2026 requires the discipline of not building until the evidence justifies it.

The work in this stage is research, customer discovery, competitive analysis, and honest evaluation of disconfirming evidence, all before asking Claude Code to generate your first line of production code.

创意阶段

Every 创业公司 创始人 starts from the same place: a problem they can't stop thinking about.

This is the 创业公司 phase where idea meets reality: 创业公司 success in 2026 requires the discipline of not building until the evidence justifies it.

The work in this stage is 研究/调研, 客户发现, 竞争分析, and honest evaluation of disconfirming evidence, all before asking Claude Code to generate your first line of 生产环境 code.

Idea stage goal

While in the Idea stage, the founder's main goal is research-oriented validation: assembling solid evidence that a real problem exists (and that your proposed solution effectively addresses it) before committing resources to building.

- Is this problem real, specific, and frequent enough to build around?
- Who exactly has it, and is that a market?
- Is anyone else solving it, and if so, how and how well?
- What would a solution actually need to do in order to solve this problem, and

does my idea do that?

The results of these inquiries add up to answer a single, ultimate question: Is this worth building?

That means getting specific before you get moving.

"People struggle with expense reporting" is an observation.

"Finance managers at mid-market companies spend four-plus hours a week reconciling submissions because their current tools don't integrate with their accounting software" is a testable hypothesis.

创意阶段目标

While in the Idea stage, the 创始人's main goal is 研究/调研-oriented 验证: assembling solid evidence that a real problem exists (and that your proposed solution effectively addresses it) before committing resources to building.

- Is this problem real, specific, and frequent enough to build around?
- Who exactly has it, and is that a market?
- Is anyone else solving it, and if so, how and how well?
- What would a solution actually need to do in order to solve this problem, and

does my idea do that?

The results of these inquiries add up to answer a single, ultimate question: Is this worth building?

That means getting specific before you get moving.

"People struggle with expense reporting" is an observation.

"Finance managers at mid-market companies spend four-plus hours a week reconciling submissions because their current tools don't integrate with their accounting software" is a testable hypothesis.

Idea stage exit criteria

The Idea stage exit condition is finding problem-solution fit. You've established qualitative evidence, primarily from real human conversations, that you're solving a real problem for real people before you start building the thing that solves it. You're ready to leave the Idea stage when you can answer yes to all three of the following: 1.

Is the problem real and specific?

Answering in the affirmative here requires that you can name exactly who experiences this problem, how often they encounter it, how severely it affects them, and what they currently do about it.

2. Does your solution address the actual problem?

Not the problem you originally assumed, but the one the validation process revealed.

Sometimes these are the same thing, but not always. 3. Do you have enough signal to justify building?

You will never have certainty at this stage, and waiting for it is its own failure mode, but you need enough qualitative evidence that committing to an MVP is a reasoned decision over an act of faith.

创意阶段退出标准

The Idea stage exit condition is finding problem-solution fit. You've established qualitative evidence, primarily from real human conversations, that you're solving a real problem for real people before you start building the thing that solves it. You're ready to leave the Idea stage when you can answer yes to all three of the following: 1.

Is the problem real and specific?

Answering in the affirmative here requires that you can name exactly who experiences this problem, how often they encounter it, how severely it affects them, and what they currently do about it.

2. Does your solution address the actual problem?

Not the problem you originally assumed, but the one the 验证 process revealed.

Sometimes these are the same thing, but not always. 3.

Do you have enough signal to justify building?

You will never have certainty at this stage, and waiting for it is its own failure mode, but you need enough qualitative evidence that committing to an 最小可行产品 (MVP) is a reasoned decision over an act of faith.

Idea stage challenges

The Idea stage is where the most important work of your startup journey happens, because it's where the most consequential mistakes are made: getting something wrong now can quickly run your budding venture right off the rails.

The majority of ideation phase challenges involve moving faster than your understanding justifies, though, so founders who proceed with thoughtfulness and deliberation will experience steady progress.

创意阶段挑战

The Idea stage is where the most important work of your 创业公司 journey happens, because it's where the most consequential mistakes are made: getting something wrong now can quickly run your budding venture right off the rails.

The majority of ideation phase challenges involve moving faster than your understanding justifies, though, so 创始人s who proceed with thoughtfulness and deliberation will experience steady progress.

Mistaking building for validating

The challenge: When technical blockers are lifted, an impassioned founder risks skipping the most important work in the startup journey: validating that their idea is genuinely a solution that people need and will use.

Even before the current era of agentic coding, 42% of startups failed because they built something nobody wanted.

Now, though, agentic coding solutions like Claude Code have drastically collapsed the distance between "I have an idea" and "I have a product" and that failure rate is only going to climb.

While there's never been a better time to be a founder with a synapse-shakingly good idea, the rapidity and ease of spinning up a prototype that looks something like a product also, counterintuitively, presents a genuinely dangerous existential risk for the AI-native startup.

Until very recently, building required real dev time and budget, and getting even a basic prototype together typically took months.

Now that the hurdle of technical development is largely gone, though, AI makes it all too easy for a founder to jump straight into building without validating its utility in the real world.

Reaching problem-solution fit requires first validating your hypothesis then building, but many first-time (and even experienced) founders mistakenly believe that AI short-circuits that requirement, turning the flow into have an idea -> immediately build a prototype -> treat the existence of the prototype as validation.

The prototype becomes a reason to believe the hypothesis was right all along, without ever testing whether it's actually true.

A working prototype is easy to mistake as concrete evidence that you're solving a real problem, but it's not.

Your prototype instead serves as a useful pressure- testing prop for conversations with potential users.

These conversations themselves are the real evidence.

Mistaking building for validating

The challenge: When technical blockers are lifted, an impassioned 创始人 risks skipping the most important work in the 创业公司 journey: validating that their idea is genuinely a solution that people need and will use.

Even before the current era of 代理编程, 42% of 创业公司s failed because they built something nobody wanted.

Now, though, 代理编程 solutions like Claude Code have drastically collapsed the distance between "I have an idea" and "I have a product" and that failure rate is only going to climb.

While there's never been a better time to be a 创始人 with a synapse-shakingly good idea, the rapidity and ease of spinning up a prototype that looks something like a product also, counterintuitively, presents a genuinely dangerous existential risk for the AI原生 创业公司.

Until very recently, building required real dev time and budget, and getting even a basic prototype together typically took months.

Now that the hurdle of technical development is largely gone, though, AI makes it all too easy for a 创始人 to jump straight into building without validating its utility in the real world.

Reaching problem-solution fit requires first validating your hypothesis then building, but many first-time (and even experienced) 创始人s mistakenly believe that AI short-circuits that requirement, turning the flow into have an idea -> immediately build a prototype -> treat the existence of the prototype as 验证.

The prototype becomes a reason to believe the hypothesis was right all along, without ever testing whether it's actually true.

A working prototype is easy to mistake as concrete evidence that you're solving a real problem, but it's not.

Your prototype instead serves as a useful pressure- testing prop for conversations with potential users.

These conversations themselves are the real evidence.

Premature scaling

The challenge: When building is effortless and instant, you can scale execution far ahead of what business demands.

Scaling prematurely means committing to a product path before you've genuinely validated that the path is worth committing to.

This has always been a startup killer, but AI has made it dramatically easier for founders to fall into the premature scaling trap without noticing.

Agentic coding assistants are so powerful that it's easy to scale execution far ahead of validating problem-solution fit without ever consciously deciding to stray off course.

It will generate, test, debug, and refactor a codebase around a fundamentally flawed premise with exactly the same enthusiasm it brings to a great idea.

The intelligence in the system is yours.

The prime directive at this stage is keeping your sense-making ahead of your building, especially when building is so quick and feels so effortless.

Premature scaling

The challenge: When building is effortless and instant, you can 规模化 执行 far ahead of what business demands.

Scaling prematurely means committing to a product path before you've genuinely validated that the path is worth committing to.

This has always been a 创业公司 killer, but AI has made it dramatically easier for 创始人s to fall into the premature scaling trap without noticing.

代理编程 assistants are so powerful that it's easy to 规模化 执行 far ahead of validating problem-solution fit without ever consciously deciding to stray off course.

It will generate, test, 调试, and refactor a codebase around a fundamentally flawed premise with exactly the same enthusiasm it brings to a great idea.

The intelligence in the system is yours.

The prime directive at this stage is keeping your sense-making ahead of your building, especially when building is so quick and feels so effortless.

Loss of objectivity

The challenge: Ask an AI tool for evidence supporting what you already believe, and it will find it.

Confirmation bias now comes with a research engine.

Confirmation bias has always been an occupational hazard in startups: founders are, by nature, passionate about their ideas.

Now, AI tools have given confirmation bias a significant powerup.

Ask AI to validate your startup idea and it will find supporting evidence; ask it to size your potential market and it will find the number that makes your TAM look fundable.

AI follows your direction, which means a founder who isn't asking hard questions can now construct an elaborate, well-researched-looking case for a bad idea faster than ever before, while feeling fully confident that they are, in fact, performing due diligence.

The antidote is the same tool, only pointed in the opposite direction: AI will pressure-test an idea just as thoroughly as it validates one.

When research and structured adversarial thinking surface evidence that your idea needs revision, this the signal to pivot.

How Claude can help Idea stage founders Progressing your AI-native startup concept through the Idea stage can feel like it takes forever.

You are a founder and you just want to build.

But this all-important kickoff phase is fundamentally a research and validation exercise, which means reaching for the tools that help you think more rigorously before going all in on writing code.

Here are ways to use Claude across its product surfaces (Chat, Claude Cowork, and Claude Code) for moving through the Idea stage as quickly as possible while doing proper due diligence.

Chat, Claude Cowork, or Claude Code: choosing the right Claude surface AI makes it easier for startup founders to ship faster, automate tedious workflows, and operate at scale, but the surface you use matters.

Here's when to use Chat, Claude Cowork, or Claude Code depending on the task at hand.

Chat is for quick exchanges without leaving the app you're already in.

Use it for the constant small tasks of running a company: pulling the one-sentence takeaway from a dense investor memo, sanity-checking a claim before a board meeting, or making sense of a long Slack thread with your team.

Claude Cowork is for the knowledge work that actually takes time: pulling from many sources, making sense of it, and producing something finished, like a doc, deck, or spreadsheet.

Loss of objectivity

The challenge: Ask an AI tool for evidence supporting what you already believe, and it will find it.

Confirmation bias now comes with a 研究/调研 engine.

Confirmation bias has always been an occupational hazard in 创业公司s: 创始人s are, by nature, passionate about their ideas.

Now, AI tools have given confirmation bias a significant powerup.

Ask AI to validate your 创业公司 idea and it will find supporting evidence; ask it to size your potential market and it will find the number that makes your TAM look fundable.

AI follows your direction, which means a 创始人 who isn't asking hard questions can now construct an elaborate, well-研究/调研ed-looking case for a bad idea faster than ever before, while feeling fully confident that they are, in fact, performing due diligence.

The antidote is the same tool, only pointed in the opposite direction: AI will pressure-test an idea just as thoroughly as it validates one.

When 研究/调研 and structured adversarial thinking surface evidence that your idea needs revision, this the signal to pivot.

How Claude can help Idea stage 创始人s Progressing your AI原生 创业公司 concept through the Idea stage can feel like it takes forever.

You are a 创始人 and you just want to build.

But this all-important kickoff phase is fundamentally a 研究/调研 and 验证 exercise, which means reaching for the tools that help you think more rigorously before going all in on writing code.

Here are ways to use Claude across its product surfaces (Chat, Claude Cowork, and Claude Code) for moving through the Idea stage as quickly as possible while doing proper due diligence.

Chat, Claude Cowork, or Claude Code: choosing the right Claude surface AI makes it easier for 创业公司 创始人s to ship faster, automate tedious workflows, and operate at 规模化, but the surface you use matters.

Here's when to use Chat, Claude Cowork, or Claude Code depending on the task at hand.

Chat is for quick exchanges without leaving the app you're already in.

Use it for the constant small tasks of running a company: pulling the one-sentence takeaway from a dense investor memo, sanity-checking a claim before a board meeting, or making sense of a long Slack thread with your team.

Claude Cowork is for the knowledge work that actually takes time: pulling from many sources, making sense of it, and producing something finished, like a doc, deck, or spreadsheet.

Think turning a folder of customer call transcripts into a themed findings doc for your next product review, building a competitive landscape from a dozen vendor sites before a fundraiser, or a standing Monday-morning task that pulls metrics from your connected tools and drops a weekly KPI brief into a shared folder.

Claude Code is the agentic coding environment for the engineers on your team: direct codebase access, Plan Mode, git integration, and local, IDE, or sandboxed cloud environments.

It's where a lean team ships features across a growing codebase, migrates legacy code from the MVP days, and moves from prototype to production without waiting on more headcount.

If the task is...

Reach for Why A question, a rewrite, a quick brainstorm Chat Fast, conversational, no setup Research, analysis, or a finished document built from your files and systems Claude Cowork Folder access, connectors, skills, scheduled runs Writing, testing, or shipping software Claude Code Codebase access, diffs, git, dev environments The three share the same Claude underneath; what changes is the workspace around it.

Think turning a folder of customer call transcripts into a themed findings doc for your next product review, building a competitive landscape from a dozen vendor sites before a fundraiser, or a standing Monday-morning task that pulls 指标 from your connected tools and drops a weekly KPI brief into a shared folder.

Claude Code is the 代理编程 environment for the engineers on your team: direct codebase access, Plan Mode, git 集成, and local, IDE, or sandboxed cloud environments.

It's where a lean team ships features across a growing codebase, migrates legacy code from the 最小可行产品 (MVP) days, and moves from prototype to 生产环境 without waiting on more 团队规模.

If the task is...

Reach for Why A question, a rewrite, a quick brainstorm Chat Fast, conversational, no setup Research, analysis, or a finished document built from your files and systems Claude Cowork Folder access, connectors, skills, scheduled runs Writing, testing, or shipping software Claude Code Codebase access, diffs, git, dev environments The three share the same Claude underneath; what changes is the workspace around it.

Defining and pressure-testing the problem hypothesis

Your own domain expertise and up-front research have already generated a hypothesis.

The first job is to sharpen it until it's actually testable.

Claude is particularly useful here for forcing specificity: who exactly has this problem, how often, how severely, and what do they currently do about it?

A problem statement that can't answer those questions precisely isn't ready to validate.

- Exercise: Work with Claude to sharpen your problem statement until it's

a testable hypothesis. For example, "Contract review takes too long" is not meaningfully testable.

But "In-house legal teams at mid-market companies spend 3+ days per contract review cycle because redlines are managed across email threads rather than a single version-controlled document" is very testable.

Your next move is to ask Claude to argue against your idea, and to find disconfirming evidence that refutes your hypothesis.

This can surface negative market signals, failed competitors, customer behavior patterns, and structural obstacles that a supportive synthesis would have quietly deprioritized.

The goal is to arrive at customer discovery having already stress-tested your assumptions against the strongest available counterarguments so that informational user interviews are genuinely open-ended rather than a search for confirmation.

Note: Using Claude as structured devil's advocate is a core use case at every stage of the AI startup life cycle.

Defining and pressure-testing the problem hypothesis

Your own domain expertise and up-front 研究/调研 have already generated a hypothesis.

The first job is to sharpen it until it's actually testable.

Claude is particularly useful here for forcing specificity: who exactly has this problem, how often, how severely, and what do they currently do about it?

A problem statement that can't answer those questions precisely isn't ready to validate.

- Exercise: Work with Claude to sharpen your problem statement until it's

a testable hypothesis.

For example, "Contract review takes too long" is not meaningfully testable.

But "In-house legal teams at mid-market companies spend 3+ days per contract review cycle because redlines are managed across email threads rather than a single version-controlled document" is very testable.

Your next move is to ask Claude to argue against your idea, and to find disconfirming evidence that refutes your hypothesis.

This can surface negative market signals, failed competitors, customer behavior patterns, and structural obstacles that a supportive synthesis would have quietly deprioritized.

The goal is to arrive at 客户发现 having already stress-tested your assumptions against the strongest available counterarguments so that informational user interviews are genuinely open-ended rather than a search for confirmation.

Note: Using Claude as structured devil's advocate is a core 用例 at every stage of the AI 创业公司 life cycle.

Sizing up your competitors

There's a startup-specific phenomenon called competitor neglect: the tendency to focus so intensely on your own vision and execution that you systematically underweight what others are doing in the same space.

Fortunately, AI offers the antidote: ask Claude to make the most compelling argument for why a competitor in this solution space would succeed while you do not.

Claude can analyze why their approach is actually better, why customers would choose them, why your potential differentiators may not be as defensible as you think.

- Exercise: Ask Claude to map your competitive landscape by tier: direct

competitors, indirect competitors, potential acquirers, and adjacent players who could move into your space.

Then ask it to argue for why each tier poses a genuine threat to your success, not just the version of the threat that's easiest to dismiss.

Sizing up your competitors

There's a 创业公司-specific phenomenon called competitor neglect: the tendency to focus so intensely on your own vision and 执行 that you systematically underweight what others are doing in the same space.

Fortunately, AI offers the antidote: ask Claude to make the most compelling argument for why a competitor in this solution space would succeed while you do not.

Claude can analyze why their approach is actually better, why customers would choose them, why your potential differentiators may not be as defensible as you think.

- Exercise: Ask Claude to map your competitive landscape by tier: direct

competitors, indirect competitors, potential acquirers, and adjacent players who could move into your space.

Then ask it to argue for why each tier poses a genuine threat to your success, not just the version of the threat that's easiest to dismiss.

Market research

Claude Code can synthesize publicly available customer feedback to surface recurring complaints and unmet needs.

Bonus: doing this is essentially free qualitative research on your competitors' customers.

- Exercise: Direct Claude Cowork to synthesize competitor reviews across your

key sources and identify the top complaints that existing solutions haven't resolved.

If your hypothesis addresses one or more of them, that's strong evidence of problem-solution fit.

If it doesn't, that's worth knowing too.

Claude Cowork can also extract relevant information and figures from dense industry reports, analyst filings, and market research documents; next, these clean, synthesized inputs become ideal context for Claude's analysis work.

- Exercise: Build TAM/SAM/SOM models from publicly available data and

pressure-test the assumptions behind them.

Identify whether the market is expanding, consolidating, or mature; this context influences how you think about timing and differentiation.

Map the buyer landscape: who holds budget, who influences decisions, and whether those are the same person.

Market 研究/调研

Claude Code can synthesize publicly available customer feedback to surface recurring complaints and unmet needs.

Bonus: doing this is essentially free qualitative 研究/调研 on your competitors' customers.

- Exercise: Direct Claude Cowork to synthesize competitor reviews across your

key sources and identify the top complaints that existing solutions haven't resolved.

If your hypothesis addresses one or more of them, that's strong evidence of problem-solution fit.

If it doesn't, that's worth knowing too.

Claude Cowork can also extract relevant information and figures from dense industry reports, analyst filings, and market 研究/调研 documents; next, these clean, synthesized inputs become ideal context for Claude's analysis work.

- Exercise: Build TAM/SAM/SOM models from publicly available data and

pressure-test the assumptions behind them.

Identify whether the market is expanding, consolidating, or mature; this context influences how you think about timing and differentiation.

Map the buyer landscape: who holds budget, who influences decisions, and whether those are the same person.

Trend analysis

Finally, use Claude to listen for early indicators that tell you whether you're entering at the right moment.

Track subreddits and LinkedIn groups where conversations about your problem are already happening and the exact language users reach for when describing their issues.

Ask Claude to identify analogous markets where a similar problem was solved, and extract what worked and what didn't.

Surface regulatory, technological, or demographic trends that could accelerate or threaten the opportunity.

- Exercise: Ask Claude to identify three external trends—regulatory,

technological, or demographic—that could significantly affect your market in the next two years, and to assess whether each one is a tailwind or a headwind for your specific hypothesis.

Note: The market research and competitive mapping work in this section isn't a one-time exercise.

You are going to continue making discoveries and evolving your thinking through the MVP and Launch stages, so it's important to repeat these exercises whenever your hypothesis evolves.

Trend analysis

Finally, use Claude to listen for early indicators that tell you whether you're entering at the right moment.

Track subreddits and LinkedIn groups where conversations about your problem are already happening and the exact language users reach for when describing their issues.

Ask Claude to identify analogous markets where a similar problem was solved, and extract what worked and what didn't.

Surface regulatory, technological, or demographic trends that could accelerate or threaten the opportunity.

- Exercise: Ask Claude to identify three external trends—regulatory,

technological, or demographic—that could significantly affect your market in the next two years, and to assess whether each one is a tailwind or a headwind for your specific hypothesis.

Note: The market 研究/调研 and competitive mapping work in this section isn't a one-time exercise.

You are going to continue making discoveries and evolving your thinking through the 最小可行产品 (MVP) and 发布阶段s, so it's important to repeat these exercises whenever your hypothesis evolves.

Plan and design customer discovery

The quality of what you learn by talking to potential users for your product is determined by (1) the quality of the questions you ask and (2) whether you are posing these to the right people.

Claude is particularly helpful for conducting customer discovery, including who to talk to, what to ask, and how to make sense of what you hear.

Plan and design 客户发现

The quality of what you learn by talking to potential users for your product is determined by (1) the quality of the questions you ask and (2) whether you are posing these to the right people.

Claude is particularly helpful for conducting 客户发现, including who to talk to, what to ask, and how to make sense of what you hear.

Who to talk to

A precise target profile is infinitely more valuable than a long contact list, including specific job titles, company types, team structures, and seniority levels most likely to experience the problem acutely.

From there, identify where those people are actually reachable—the communities, events, LinkedIn groups, and Slack workspaces where they congregate—and build a prioritization framework for who to reach out to first based on how close they are to the problem.

Who to talk to

A precise target profile is infinitely more valuable than a long contact list, including specific job titles, company types, team structures, and seniority levels most likely to experience the problem acutely.

From there, identify where those people are actually reachable—the communities, events, LinkedIn groups, and Slack workspaces where they congregate—and build a prioritization framework for who to reach out to first based on how close they are to the problem.

What to ask

With your targets defined, use Claude to build the interview framework itself: the right questions, in the right order, structured to surface what people actually do rather than what they think they would do.

A rookie founder mistake is asking a generic, open-ended question about the future ("would you use something like this?

") instead of specifically querying the relevant past ("tell me about the last time you dealt with this problem.

") Claude can flag where your draft questions are leading the respondent, too broad, or otherwise likely to generate noise instead of signal.

Claude can also help you in designing follow-up questions to probe deflections or drill down on vague answers to important questions.

If your hypothesis involves more than one persona, Claude can also design different question sets for each.

A finance manager and a CFO have different relationships to the same problem, and a single interview framework will flatten that distinction.

- Exercise: Draft your interview questions by hand first, ask Claude to audit

them.

Ask it specifically to flag any question that is leading, future-facing, too broad, or likely to produce a socially desirable answer rather than an honest one.

Then ask it to suggest a follow-up probe for the two or three moments in the interview most likely to generate deflection.

What to ask

With your targets defined, use Claude to build the interview framework itself: the right questions, in the right order, structured to surface what people actually do rather than what they think they would do.

A rookie 创始人 mistake is asking a generic, open-ended question about the future ("would you use something like this?

") instead of specifically querying the relevant past ("tell me about the last time you dealt with this problem.

") Claude can flag where your draft questions are leading the respondent, too broad, or otherwise likely to generate noise instead of signal.

Claude can also help you in designing follow-up questions to probe deflections or drill down on vague answers to important questions.

If your hypothesis involves more than one persona, Claude can also design different question sets for each.

A finance manager and a CFO have different relationships to the same problem, and a single interview framework will flatten that distinction.

- Exercise: Draft your interview questions by hand first, ask Claude to audit

them.

Ask it specifically to flag any question that is leading, future-facing, too broad, or likely to produce a socially desirable answer rather than an honest one.

Then ask it to suggest a follow-up probe for the two or three moments in the interview most likely to generate deflection.

Post-interview analysis

After each conversation, use Claude to debrief: feed it your notes and ask it to identify what confirmed your hypothesis, what challenged it, and what was genuinely surprising.

Once you've gathered a batch of interviews, run your full set of interview notes through Claude Cowork to surface recurring themes, contradictions, and the strongest signals in both directions.

Then take that synthesized output back to Claude and ask it to flag where your own read of the data might be pattern-matching to what you want to hear rather than what's actually there.

- Exercise: After every five interviews, direct Claude Cowork to synthesize your

notes and produce two lists: the evidence that supports your hypothesis, and the evidence that challenges it.

If the first list is significantly longer than the second, ask Claude whether that asymmetry reflects what's actually in the data—or what you were hoping to find.

Post-interview analysis

After each conversation, use Claude to debrief: feed it your notes and ask it to identify what confirmed your hypothesis, what challenged it, and what was genuinely surprising.

Once you've gathered a batch of interviews, run your full set of interview notes through Claude Cowork to surface recurring themes, contradictions, and the strongest signals in both directions.

Then take that synthesized output back to Claude and ask it to flag where your own read of the data might be pattern-matching to what you want to hear rather than what's actually there.

- Exercise: After every five interviews, direct Claude Cowork to synthesize your

notes and produce two lists: the evidence that supports your hypothesis, and the evidence that challenges it.

If the first list is significantly longer than the second, ask Claude whether that asymmetry reflects what's actually in the data—or what you were hoping to find.

Customer outreach and scheduling

Use Claude Cowork to automate the operational lift around building a contact list, running outreach, and scheduling user interviews.

Claude Cowork can use the target profile you defined with Claude (including job titles, company types, and seniority levels) to research and compile a structured list of prospects and verified contact information.

It then drafts personalized outreach emails at scale, tailoring each one to the individual's role and context.

As responses come in, it connects to Gmail and Google Calendar via MCP to manage the thread, handle scheduling requests, and get interviews on the calendar.

The workflow continues as Claude Cowork generates follow-up drafts on a defined cadence (a day-seven follow-up for contacts who haven't responded, for instance) and updates your tracking sheet as each step completes so you always know where every prospect stands in the pipeline.

- Exercise: Give Claude Cowork your validated interview target profile and

ask it to build a prospect list, draft a personalized outreach sequence, and set up a tracking sheet with columns for outreach status, follow-up cadence, and interview completion.

Then let it run the coordination while you focus on preparing for the conversations themselves.

Customer outreach and scheduling

Use Claude Cowork to automate the operational lift around building a contact list, running outreach, and scheduling user interviews.

Claude Cowork can use the target profile you defined with Claude (including job titles, company types, and seniority levels) to 研究/调研 and compile a structured list of prospects and verified contact information.

It then drafts personalized outreach emails at 规模化, tailoring each one to the individual's role and context.

As responses come in, it connects to Gmail and Google Calendar via MCP to manage the thread, handle scheduling requests, and get interviews on the calendar.

The workflow continues as Claude Cowork generates follow-up drafts on a defined cadence (a day-seven follow-up for contacts who haven't responded, for instance) and updates your tracking sheet as each step completes so you always know where every prospect stands in the 流水线/管道.

- Exercise: Give Claude Cowork your validated interview target profile and

ask it to build a prospect list, draft a personalized outreach sequence, and set up a tracking sheet with columns for outreach status, follow-up cadence, and interview completion.

Then let it run the coordination while you focus on preparing for the conversations themselves.

Design your final solution concept

You've done the validation work: the problem is real, you know who has it, and you have a solution concept that the evidence supports.

Use Claude to develop and challenge your solution concept from every angle: What are the gaps? What alternatives exist?

What would have to be true for this solution to work at scale?

This is an important reality checkpoint: does this design actually address the problem the validation process revealed, and not the problem you originally assumed going in?

- Exercise: Present your solution concept to Claude and ask it to identify the

three assumptions your design depends on most heavily.

Then ask what would have to be true for each assumption to hold, and what the consequences are if any one of them doesn't.

Design your final solution concept

You've done the 验证 work: the problem is real, you know who has it, and you have a solution concept that the evidence supports.

Use Claude to develop and challenge your solution concept from every angle: What are the gaps?

What alternatives exist?

What would have to be true for this solution to work at 规模化?

This is an important reality checkpoint: does this design actually address the problem the 验证 process revealed, and not the problem you originally assumed going in?

- Exercise: Present your solution concept to Claude and ask it to identify the

three assumptions your design depends on most heavily.

Then ask what would have to be true for each assumption to hold, and what the consequences are if any one of them doesn't.

Build a lightweight prototype with Claude Code

Now for the fun part: with a validated hypothesis and a stress-tested solution concept, you're finally ready to build something.

This is the moment in the Idea stage where Claude Code enters the picture.

Even if you've been tinkering all along, now is the point where you generate your official lightweight prototype: the minimum surface area needed to put your idea in front of a real human and get a genuine reaction.

You're not building a real-world product (yet); you're constructing a functional sample of your idea to use in customer and investor conversations.

Real users reacting to something they can actually touch will tell you things that a dozen problem-solution discovery interviews couldn't.

Before, you were establishing that the problem you're solving is real; now, you are asking potential users to engage with the proposed solution.

- Exercise: Define the single core interaction your solution depends on. Direct

Claude Code to build only that.

When you have it, put it in front of five people from your validated target profile and ask them to try it out.

What you learn in those five conversations determines whether you keep building, or go back to the drawing board.

Reaching the end of the Idea stage is a giant leap ahead in the AI startup race because now you're not betting on a hunch; you're executing against evidence.

Now comes the MVP stage, where the founder's guiding question goes from "Is this worth building?"

" to "What exactly should we build first?" and AI's primary role shifts from research partner to construction crew.

Build a lightweight prototype with Claude Code

Now for the fun part: with a validated hypothesis and a stress-tested solution concept, you're finally ready to build something.

This is the moment in the Idea stage where Claude Code enters the picture.

Even if you've been tinkering all along, now is the point where you generate your official lightweight prototype: the minimum surface area needed to put your idea in front of a real human and get a genuine reaction.

You're not building a real-world product (yet); you're constructing a functional sample of your idea to use in customer and investor conversations.

Real users reacting to something they can actually touch will tell you things that a dozen problem-solution discovery interviews couldn't.

Before, you were establishing that the problem you're solving is real; now, you are asking potential users to engage with the proposed solution.

- Exercise: Define the single core interaction your solution depends on.

Claude Code to build only that.

When you have it, put it in front of five people from your validated target profile and ask them to try it out.

What you learn in those five conversations determines whether you keep building, or go back to the drawing board.

Reaching the end of the Idea stage is a giant leap ahead in the AI 创业公司 race because now you're not betting on a hunch; you're executing against evidence.

Now comes the 最小可行产品 (MVP) stage, where the 创始人's guiding question goes from "Is this worth building?"

" to "What exactly should we build first?"

" and AI's primary role shifts from 研究/调研 partner to construction crew.

MVP Stage

Plenty of founders treat the MVP stage as a construction phase, but the MVP stage is still fundamentally an evidence-gathering exercise.

The difference is that you are now gathering evidence about the solution instead of the problem space; specifically, whether a real, identifiable group of people finds it valuable enough to use it, return to it, pay for it, and/or tell others about it.

MVP stage goals As the founder of an AI-native startup, your goal is to translate a validated problem into a working product that real users will actually use.

This is not the full version with every roadmap feature but the smallest, most focused iteration of your idea that puts a real solution in front of real users and generates real evidence of product-market fit.

At the same time, how you build now determines what's possible later.

This means that the MVP stage has a second, equally important goal of moving fast without accruing the type of technical debt that compounds—and will haunt you the moment real users arrive in meaningful numbers.

And finally, investing in persistent context from day one is what keeps AI a force multiplier instead of a source of entropy.

In an AI-native startup, your codebase is something you collaborate with AI on session after session, which makes legibility foundational.

Founders who skip specs, architectural decisions, and context files (like CLAUDE.

md) hit a predictable wall where every new session requires re-explaining the codebase and AI-generated changes drift from the original vision.

最小可行产品阶段

Plenty of 创始人s treat the 最小可行产品 (MVP) stage as a construction phase, but the 最小可行产品 (MVP) stage is still fundamentally an evidence-gathering exercise.

The difference is that you are now gathering evidence about the solution instead of the problem space; specifically, whether a real, identifiable group of people finds it valuable enough to use it, return to it, pay for it, and/or tell others about it.

最小可行产品阶段目标s As the 创始人 of an AI原生 创业公司, your goal is to translate a validated problem into a working product that real users will actually use.

This is not the full version with every 路线图 feature but the smallest, most focused 迭代 of your idea that puts a real solution in front of real users and generates real evidence of 产品-市场契合.

At the same time, how you build now determines what's possible later.

This means that the 最小可行产品 (MVP) stage has a second, equally important goal of moving fast without accruing the type of technical debt that compounds—and will haunt you the moment real users arrive in meaningful numbers.

And finally, investing in persistent context from day one is what keeps AI a force multiplier instead of a source of entropy.

In an AI原生 创业公司, your codebase is something you collaborate with AI on session after session, which makes legibility foundational.

Founders who skip specs, architectural decisions, and context files (like CLAUDE.

md) hit a predictable wall where every new session requires re-explaining the codebase and AI-generated changes drift from the original vision.

MVP stage exit criteria

The MVP stage exit condition is genuine evidence of product-market fit: proof that a specific, identifiable group of users has found the product valuable enough to return to it (retention), pay for it (revenue), or tell others about it (referral).

最小可行产品阶段退出标准

The 最小可行产品 (MVP) stage exit condition is genuine evidence of 产品-市场契合: proof that a specific, identifiable group of users has found the product valuable enough to return to it (retention), pay for it (营收), or tell others about it (referral).

MVP stage challenges

In the MVP stage, a founder's prime directives are speed and judgment.

The challenges here center on whether you can build the right thing, the right way, fast enough to matter, without cutting corners that will cost you later.

最小可行产品阶段挑战

In the 最小可行产品 (MVP) stage, a 创始人's prime directives are speed and judgment.

The challenges here center on whether you can build the right thing, the right way, fast enough to matter, without cutting corners that will cost you later.

Agentic technical debt

The challenge: Because AI essentially removes every natural bottleneck that once controlled what reaches production, speed is guaranteed.

But when speed is the only variable that founders factor into their MVP build, they risk accruing technical debt they'll struggle to pay off.

Some technical debt is appropriate at the MVP stage, with the understanding that it must be managed before scaling.

It builds gradually and can be cleared over time or in a dedicated sprint. AI technical debt, however, compounds.

Without specs and architectural constraints written down somewhere the AI can read, each session re-derives foundational decisions from scratch, and those decisions drift.

You end up with a codebase that has no coherent mental model behind it, not because any single piece is bad, but because the pieces were never designed to fit together.

That's a real problem, and it does tend to surface late.

Agentic technical debt

The challenge: Because AI essentially removes every natural bottleneck that once controlled what reaches 生产环境, speed is guaranteed.

But when speed is the only variable that 创始人s factor into their 最小可行产品 (MVP) build, they risk accruing technical debt they'll struggle to pay off.

Some technical debt is appropriate at the 最小可行产品 (MVP) stage, with the understanding that it must be managed before scaling.

It builds gradually and can be cleared over time or in a dedicated sprint.

AI technical debt, however, compounds.

Without specs and architectural constraints written down somewhere the AI can read, each session re-derives foundational decisions from scratch, and those decisions drift.

You end up with a codebase that has no coherent mental model behind it, not because any single piece is bad, but because the pieces were never designed to fit together.

That's a real problem, and it does tend to surface late.

Falling for false product-market fit

The challenge: AI tools can generate impressive early numbers, but these are not a guarantee that the market needs your product.

Early momentum is one of the most psychologically powerful experiences a founder can have.

After weeks or months of validation work and careful, disciplined building, shipping a product feels like confirmation that you were right all along.

Agentic coding tools can help you reach this moment faster than ever before, but early traction is not the same as product-market fit.

Launch energy is generated from ephemeral forces, like your founder's friends, prospective buyers at your investor's other portfolio companies, or a Hacker News headline that drives a spike.

Unfortunately, none of these reliably predicts what happens at week six or week twelve when that initial boost has faded.

Falling for false 产品-市场契合

The challenge: AI tools can generate impressive early numbers, but these are not a guarantee that the market needs your product.

Early momentum is one of the most psychologically powerful experiences a 创始人 can have.

After weeks or months of 验证 work and careful, disciplined building, shipping a product feels like confirmation that you were right all along.

代理编程 tools can help you reach this moment faster than ever before, but early traction is not the same as 产品-市场契合.

Launch energy is generated from ephemeral forces, like your 创始人's friends, prospective buyers at your investor's other portfolio companies, or a Hacker News headline that drives a spike.

Unfortunately, none of these reliably predicts what happens at week six or week twelve when that initial boost has faded.

Zero-friction scope creep

The challenge: When building feels effortless and is nearly free, there's always one more cool feature to add or one more edge case to handle.

This scope creep can do more harm than good. Scope creep has always been a startup risk.

The difference now is that the traditional forcing function against it—the real cost of engineering time—no longer exists in the same way when adding a feature takes an afternoon instead of a sprint.

The challenge here is that each individual addition is defensible.

Of course the product should handle that edge case; of course users will want that workflow.

These don't feel like scope creep in the moment because each one takes so little effort to build with agentic coding, but as your product sprawls beyond its original boundaries you risk losing direction and momentum.

The antidote is a written scope definition created before building begins, describing what the product does, what it deliberately does not do, and the specific evidence from real users that would justify adding something new.

This moves the decision point from "should we build this?"

" to "a critical mass of users have told us they can't get value from the product without this?"

Zero-friction scope creep

The challenge: When building feels effortless and is nearly free, there's always one more cool feature to add or one more 边界情况 to handle.

This scope creep can do more harm than good.

Scope creep has always been a 创业公司 risk.

The difference now is that the traditional forcing function against it—the real cost of engineering time—no longer exists in the same way when adding a feature takes an afternoon instead of a sprint.

The challenge here is that each individual addition is defensible.

Of course the product should handle that 边界情况; of course users will want that workflow.

These don't feel like scope creep in the moment because each one takes so little effort to build with 代理编程, but as your product sprawls beyond its original boundaries you risk losing direction and momentum.

The antidote is a written scope definition created before building begins, describing what the product does, what it deliberately does not do, and the specific evidence from real users that would justify adding something new.

This moves the decision point from "should we build this?"

" to "a critical mass of users have told us they can't get value from the product without this?"

Insecure by inexperience

The challenge: Founders using AI tools to rush applications to market without first understanding fundamental security principles end up exposing their users to preventable risks.

The hard truth is that agentic coding tools generate code that works, not code that is inherently secure.

Functional code is easy, because either the feature works or it doesn't.

Security vulnerabilities are invisible until they're exploited, which means there's no natural feedback loop to alert a first-time founder that something is wrong.

Shipping a live MVP to real users, however, means real data, real exposure, and real consequences if something goes wrong.

Slighting security isn't brand new to AI-native projects.

Bootstrapped startups in every era often delayed security considerations until late in the build, sometimes waiting until the verge of production launch.

A security review before any user touches your app or solution is the minimum responsible threshold for releasing a minimum viable product into the world.

Insecure by inexperience

The challenge: Founders using AI tools to rush applications to market without first understanding fundamental security principles end up exposing their users to preventable risks.

The hard truth is that 代理编程 tools generate code that works, not code that is inherently secure.

Functional code is easy, because either the feature works or it doesn't.

Security vulnerabilities are invisible until they're exploited, which means there's no natural 反馈循环 to alert a first-time 创始人 that something is wrong.

Shipping a live 最小可行产品 (MVP) to real users, however, means real data, real exposure, and real consequences if something goes wrong.

Slighting security isn't brand new to AI原生 projects.

Bootstrapped 创业公司s in every era often delayed security considerations until late in the build, sometimes waiting until the verge of 生产环境 发布/上线.

A security review before any user touches your app or solution is the minimum responsible threshold for releasing a minimum viable product into the world.

Define your architecture before you build

Before Claude Code writes a line of production code, use Claude to define and document the architectural decisions that will govern everything built in this stage: the patterns to follow, the dependencies to avoid, the tradeoffs being made and why.

This output will serve as a focused architectural context document and establish the guardrails that Claude Code will operate inside.

Without this context, each session starts from scratch and Claude Code is forced to infer its own structural assumptions.

Letting Claude Code build without guardrails produces a codebase that will be functional but structurally incoherent, and iterating on and scaling incoherent codebases is ultimately a waste of time and tokens.

Sooner or later there's a point where the code inevitably collapses, forcing you to rebuild from scratch.

- Exercise: Before opening Claude Code, open Claude and describe what

you're building: the core problem it solves, the users it serves, and the scale you realistically expect in the next six months.

Ask it to help you define the architectural principles that should govern your MVP build, the dependencies to avoid given your constraints, and the tradeoffs you're consciously accepting at this stage.

Next, save this output as CLAUDE.md markdown file(s).

This is your architectural context document: the first artifact of your build, and the one every subsequent session depends on.

CLAUDE.

md files serve as project- level instructions for Claude Code, providing project-specific context and instructions that are automatically read by the Agent SDK when it runs in a directory.

Functionally, they are persistent "memory" for your project.

Define your architecture before you build

Before Claude Code writes a line of 生产环境 code, use Claude to define and document the architectural decisions that will govern everything built in this stage: the patterns to follow, the dependencies to avoid, the tradeoffs being made and why.

This output will serve as a focused architectural context document and establish the guardrails that Claude Code will operate inside.

Without this context, each session starts from scratch and Claude Code is forced to infer its own structural assumptions.

Letting Claude Code build without guardrails produces a codebase that will be functional but structurally incoherent, and iterating on and scaling incoherent codebases is ultimately a waste of time and tokens.

Sooner or later there's a point where the code inevitably collapses, forcing you to rebuild from scratch.

- Exercise: Before opening Claude Code, open Claude and describe what

you're building: the core problem it solves, the users it serves, and the 规模化 you realistically expect in the next six months.

Ask it to help you define the architectural principles that should govern your 最小可行产品 (MVP) build, the dependencies to avoid given your constraints, and the tradeoffs you're consciously accepting at this stage.

Next, save this output as CLAUDE.md markdown file(s).

This is your architectural context document: the first artifact of your build, and the one every subsequent session depends on.

CLAUDE.

md files serve as project- level instructions for Claude Code, providing project-specific context and instructions that are automatically read by the Agent SDK when it runs in a directory.

Functionally, they are persistent "memory" for your project.

Define and enforce your MVP scope

Scope creep without friction is one of the defining failure modes of AI-era MVPs.

Just as you defined and documented your product's application architecture, you also need to define your MVP's scope before a single feature gets built.

Claude can help you create a scope document that describes what your MVP product does, what it deliberately does not do, and feature amendment criteria: what specific evidence from real users would justify adding something new at this point.

When new feature ideas surface—and they surely will—you use Claude to pressure-test whether it's genuine signal from users or founder enthusiasm dressed up as product thinking.

Define and enforce your 最小可行产品 (MVP) scope

Scope creep without friction is one of the defining failure modes of AI-era 最小可行产品 (MVP)s.

Just as you defined and documented your product's application architecture, you also need to define your 最小可行产品 (MVP)'s scope before a single feature gets built.

Claude can help you create a scope document that describes what your 最小可行产品 (MVP) product does, what it deliberately does not do, and feature amendment criteria: what specific evidence from real users would justify adding something new at this point.

When new feature ideas surface—and they surely will—you use Claude to pressure-test whether it's genuine signal from users or 创始人 enthusiasm dressed up as product thinking.

Build your MVP with Claude Code

Once architecture and scope are defined, Claude Code becomes the primary MVP build tool.

Use it to generate, test, debug, and iterate on your codebase, but treat each session as an execution of product decisions you have already made, not as an opportunity to throw in some new ones.

Start each Claude Code session by (1) revisiting your scope document and (2) providing the model with your CLAUDE.

md architectural context document. End each session by updating it with any decisions the session surfaced.

The goal is a codebase whose structure you can explain, not just a codebase that runs.

- Exercise: Create a simple session template for your Claude Code work that

includes the architectural context document, the specific task for this session, and any constraints or patterns to observe.

At the end of each session, add a brief log entry to the context document that details what was built, what decisions were made, and what assumptions the session introduced.

Five minutes of documentation per session is cheap insurance against architectural drift that compounds into an unmanageable codebase.

Build your 最小可行产品 (MVP) with Claude Code

Once architecture and scope are defined, Claude Code becomes the primary 最小可行产品 (MVP) build tool.

Use it to generate, test, 调试, and iterate on your codebase, but treat each session as an 执行 of product decisions you have already made, not as an opportunity to throw in some new ones.

Start each Claude Code session by (1) revisiting your scope document and (2) providing the model with your CLAUDE.

md architectural context document.

End each session by updating it with any decisions the session surfaced.

The goal is a codebase whose structure you can explain, not just a codebase that runs.

- Exercise: Create a simple session template for your Claude Code work that

includes the architectural context document, the specific task for this session, and any constraints or patterns to observe.

At the end of each session, add a brief log entry to the context document that details what was built, what decisions were made, and what assumptions the session introduced.

Five minutes of documentation per session is cheap insurance against architectural drift that compounds into an unmanageable codebase.

Security review before any user touches it

As an AI-native startup founder, your responsibility is to know what's in your codebase, understand any potential exposure vectors, and not ship obvious vulnerabilities to real users who are trusting you with their data.

Claude can do a useful first-pass security review of AI-generated code and can help identify common vulnerabilities.

It's a good habit to build into the loop before shipping.

It is not a substitute for security tooling, however, or, at higher stakes, a human reviewer—and founders who treat it as one are the ones who end up in the breach stories.

Claude Code Security goes further: it scans codebases for security vulnerabilities and suggests targeted patches for human review, surfacing issues that traditional methods can miss.

Note: At the time of this ebook's publication, Claude Code Security is a limited beta release so check current availability before building it into your workflow.

- Exercise: Before deploying to any real users, run your core application code

through Claude with a specific brief: review for authentication and session handling, data exposure in API responses, input validation and injection risks, and dependencies with known vulnerabilities.

Treat each finding seriously and assess whether it requires a fix, with human review for anything that touches authentication, secrets, or data handling.

Security review before any user touches it

As an AI原生 创业公司 创始人, your responsibility is to know what's in your codebase, understand any potential exposure vectors, and not ship obvious vulnerabilities to real users who are trusting you with their data.

Claude can do a useful first-pass security review of AI-generated code and can help identify common vulnerabilities.

It's a good habit to build into the loop before shipping.

It is not a substitute for security tooling, however, or, at higher stakes, a human reviewer—and 创始人s who treat it as one are the ones who end up in the breach stories.

Claude Code Security goes further: it scans codebases for security vulnerabilities and suggests targeted patches for human review, surfacing issues that traditional methods can miss.

Note: At the time of this ebook's publication, Claude Code Security is a limited beta release so check current availability before building it into your workflow.

- Exercise: Before 部署ing to any real users, run your core application code

through Claude with a specific brief: review for authentication and session handling, data exposure in API responses, input 验证 and injection risks, and dependencies with known vulnerabilities.

Treat each finding seriously and assess whether it requires a fix, with human review for anything that touches authentication, secrets, or data handling.

Build your measurement framework before launch

The founders who mis-identify early traction as product-market fit are typically the same ones who started tracking data after launch, using metrics chosen to assess what was working rather than to surface what wasn't.

The antidote is to establish your measurement framework before the first user shows up.

Use Claude to define which metrics matter for your specific product, what the benchmarks are, and what patterns in the data would constitute genuine product-market fit versus flattering noise.

Specifically: set your retention benchmarks, your activation criteria, and your Day 7 and Day 30 targets before releasing your MVP.

Next, define what a false positive looks like for your specific product: signups without activation, revenue without retention, or initial enthusiasm without repeat usage, for example.

When the data arrives, ask Claude to make the adversarial case against your own traction: what would a skeptic say about these numbers?

Build your measurement framework before 发布/上线

The 创始人s who mis-identify early traction as 产品-市场契合 are typically the same ones who started tracking data after 发布/上线, using 指标 chosen to assess what was working rather than to surface what wasn't.

The antidote is to establish your measurement framework before the first user shows up.

Use Claude to define which 指标 matter for your specific product, what the 基准s are, and what patterns in the data would constitute genuine 产品-市场契合 versus flattering noise.

Specifically: set your retention 基准s, your activation criteria, and your Day 7 and Day 30 targets before releasing your 最小可行产品 (MVP).

Next, define what a false positive looks like for your specific product: signups without activation, 营收 without retention, or initial enthusiasm without repeat usage, for example.

When the data arrives, ask Claude to make the adversarial case against your own traction: what would a skeptic say about these numbers?

Manage discovery and user feedback logistics

Once real users are in the product, the operational layer expands fast.

Claude Cowork handles the important-but-tedious work like building and maintaining user contact lists, running outreach sequences, scheduling feedback sessions, triaging bug reports, and tracking iteration cycles.

The same MCP integrations that managed discovery logistics in the Idea stage apply here.

Keep a human in the collection loop for nuanced exploration of user feedback.

A user saying, for example, "this is great but I wish it could also...

," requires interpretation: Is it a core need or a nice-to-have?

Is it specific to this customer or representative of a segment?

Is the missing feature the real problem, or is something upstream in the onboarding? No tool can answer those questions.

- Exercise: Configure Claude Cowork to run your MVP-stage feedback loop:

draft outreach to your early user list, schedule feedback sessions, design structured intake process for bug reports and feature requests, and write up a weekly synthesis of what's come in.

Review the synthesis yourself first; after that, you can ask Claude to analyze the information to catch any significant points you may have overlooked.

Manage discovery and user feedback logistics

Once real users are in the product, the operational layer expands fast.

Claude Cowork handles the important-but-tedious work like building and maintaining user contact lists, running outreach sequences, scheduling feedback sessions, triaging bug reports, and tracking 迭代 cycles.

The same MCP 集成s that managed discovery logistics in the Idea stage apply here.

Keep a human in the collection loop for nuanced exploration of user feedback.

A user saying, for example, "this is great but I wish it could also...

," requires interpretation: Is it a core need or a nice-to-have?

Is it specific to this customer or representative of a segment?

Is the missing feature the real problem, or is something upstream in the 入门/上手?

No tool can answer those questions.

- Exercise: Configure Claude Cowork to run your 最小可行产品 (MVP)-stage 反馈循环:

draft outreach to your early user list, schedule feedback sessions, design structured intake process for bug reports and feature requests, and write up a weekly synthesis of what's come in.

Review the synthesis yourself first; after that, you can ask Claude to analyze the information to catch any significant points you may have overlooked.

Iterate toward evidence, not toward completeness

The MVP stage ends when you have genuine evidence of product-market fit, no matter how “finished” the product feels.

Declaring that you’ve achieved product- market fit and are now ready to move on from the MVP phase to the Launch stage is ultimately a judgement exercise that combines founder intuition with collected evidence.

- The Sean Ellis test: Ask your active users: "How would you feel if you could no

longer use this product?" If more than 40% answer "very disappointed," that's a meaningful PMF indicator.

- The effort test: Pre-product-market fit, retention requires constant

intervention, including frequent outreach, incentives, personal follow-up, and a heroic founder energy expended to keep users engaged.

Post product- market fit, the product starts doing that work on its own.

When things begin pulling instead of pushing, that shift in effort is one of the clearest signals that something real has changed.

Ultimately, no single data point confirms product-market fit because it's a pattern that has to hold across multiple iteration cycles before you can definitively call it.

Iterate toward evidence, not toward completeness

The 最小可行产品 (MVP) stage ends when you have genuine evidence of 产品-市场契合, no matter how “finished” the product feels.

Declaring that you’ve achieved product- market fit and are now ready to move on from the 最小可行产品 (MVP) phase to the 发布阶段 is ultimately a judgement exercise that combines 创始人 intuition with collected evidence.

The Sean Ellis test: Ask your active users: "How would you feel if you could no longer use this product?

" If more than 40% answer "very disappointed," that's a meaningful PMF indicator.

The effort test: Pre-产品-市场契合, retention requires constant

intervention, including frequent outreach, incentives, personal follow-up, and a heroic 创始人 energy expended to keep users engaged.

Post product- market fit, the product starts doing that work on its own.

When things begin pulling instead of pushing, that shift in effort is one of the clearest signals that something real has changed.

Ultimately, no single data point confirms 产品-市场契合 because it's a pattern that has to hold across multiple 迭代 cycles before you can definitively call it.

Pivot when the evidence demands it

What if, even after investing all this work, you just can't seem to arrive at product- market fit?

The fact that your results don't confirm the direction you started from is not failure, it's the system working: the MVP stage is designed to surface this information before you over-invest in the wrong answer.

When the data doesn't support your current product, use Claude to work through what that data is telling you.

- Exploring alternative customer segments. Perhaps the users who aren't

converting were never the right target to begin with.

Often the right audience is already in your data, just underweighted.

- Adjusting your product's value prop. Maybe you have the correct audience

but your MVP is just not resonating with users.

An adjustment to onboarding, messaging, or core feature emphasis can potentially fix this without changing what you've built.

- Exercise: If you've completed three or more iteration cycles without

meaningful movement toward your product-market fit benchmarks, use Claude to run a diagnostic before deciding what to do next.

- Is there a segment in this data responding differently than the rest?
- Is the gap between designed value and experienced value a positioning

problem or a product problem?

- What would have to be true for the current product to find genuine PMF,

and is that scenario realistic given what you're seeing?

Let the answers determine whether you adjust, pivot, or return to the Idea stage.

Pivot when the evidence demands it

What if, even after investing all this work, you just can't seem to arrive at product- market fit?

The fact that your results don't confirm the direction you started from is not failure, it's the system working: the 最小可行产品 (MVP) stage is designed to surface this information before you over-invest in the wrong answer.

When the data doesn't support your current product, use Claude to work through what that data is telling you.

- Exploring alternative customer segments.

converting were never the right target to begin with.

Often the right audience is already in your data, just underweighted.

- Adjusting your product's value prop.

but your 最小可行产品 (MVP) is just not resonating with users.

An adjustment to 入门/上手, messaging, or core feature emphasis can potentially fix this without changing what you've built.

- Exercise: If you've completed three or more 迭代 cycles without

meaningful movement toward your 产品-市场契合 基准s, use Claude to run a diagnostic before deciding what to do next.

- Is there a segment in this data responding differently than the rest?

- Is the gap between designed value and experienced value a positioning

problem or a product problem?

- What would have to be true for the current product to find genuine PMF,

and is that scenario realistic given what you're seeing?

Let the answers determine whether you adjust, pivot, or return to the Idea stage.

Launch stage

If the MVP stage was about proving your product deserves to exist, the Launch stage is about proving your business deserves to grow.

Launch stage goals In the Launch stage, startup founders must turn early traction into a repeatable, sustainable growth engine.

Beyond making your product production-ready, you also must harden the infrastructure underneath it while simultaneously building an actual company around your product.

Startups are naturally founder-centric during the Idea and MVP stages because you need the full situational awareness and tight feedback loops.

Now, though, founders who still try to personally hold every thread become a Launch stage bottleneck.

The goal isn't to remove yourself from the company, but to build operational systems that free your attention for the decisions only a founder can make.

发布阶段

If the 最小可行产品 (MVP) stage was about proving your product deserves to exist, the 发布阶段 is about proving your business deserves to grow.

发布阶段目标s In the 发布阶段, 创业公司 创始人s must turn early traction into a repeatable, sustainable growth engine.

Beyond making your product 生产环境-ready, you also must harden the 基础设施 underneath it while simultaneously building an actual company around your product.

Startups are naturally 创始人-centric during the Idea and 最小可行产品 (MVP) stages because you need the full situational awareness and tight 反馈循环s.

Now, though, 创始人s who still try to personally hold every thread become a 发布阶段 bottleneck.

The goal isn't to remove yourself from the company, but to build operational systems that free your attention for the decisions only a 创始人 can make.

Launch stage exit criteria

The Launch stage exit condition has three elements: 1. Growth is repeatable and channel-driven.

You're not just retaining users, you're acquiring them predictably through specific channels with understood unit economics: CAC, LTV, and payback period are numbers you know and can defend.

2. The product can handle production workloads.

Infrastructure is hardened, security and compliance are in order, and reliability holds under real production conditions (not just the conditions you tested for).

3. Operations run without founder bottlenecks. Processes exist and automation is in place.

You are no longer the person personally handling support, triage, sprint planning, or reporting.

发布阶段退出标准

The 发布阶段 exit condition has three elements: 1.

Growth is repeatable and channel-driven.

You're not just retaining users, you're acquiring them predictably through specific channels with understood unit economics: CAC, LTV, and payback period are numbers you know and can defend.

2. The product can handle 生产环境 workloads.

Infrastructure is hardened, security and compliance are in order, and reliability holds under real 生产环境 conditions (not just the conditions you tested for).

3. Operations run without 创始人 bottlenecks.

Processes exist and automation is in place.

You are no longer the person personally handling support, triage, sprint planning, or reporting.

Launch stage challenges

Finding product-market fit is the hardest problem in the early startup lifecycle.

Now, the founder's challenge becomes keeping it.

The Launch stage is where companies that found real product traction may still fall apart if the organization that surrounds and supports the product can't keep up.

These are the failure modes to watch for.

发布阶段挑战

Finding 产品-市场契合 is the hardest problem in the early 创业公司 lifecycle.

Now, the 创始人's challenge becomes keeping it.

The 发布阶段 is where companies that found real product traction may still fall apart if the organization that surrounds and supports the product can't keep up.

These are the failure modes to watch for.

Technical debt comes due

The challenge: The MVP codebase built for speed and validation ran well enough to prove the product worked, but production traffic, new features, and growing complexity are now exposing the shortcuts.

At MVP, accumulating some technical debt was a reasonable tradeoff for velocity.

In the Launch phase, that debt starts accruing interest, and the longer it goes unaddressed, the more expensive it is to fix.

The solution consists of a systematic architectural audit to identify structural weaknesses, targeted refactoring to address the worst of them, and a meaningful expansion of test coverage so that the next round of feature work doesn't reintroduce the same problems.

Technical debt comes due

The challenge: The 最小可行产品 (MVP) codebase built for speed and 验证 ran well enough to prove the product worked, but 生产环境 traffic, new features, and growing complexity are now exposing the shortcuts.

At 最小可行产品 (MVP), accumulating some technical debt was a reasonable tradeoff for velocity.

In the Launch phase, that debt starts accruing interest, and the longer it goes unaddressed, the more expensive it is to fix.

The solution consists of a systematic architectural audit to identify structural weaknesses, targeted refactoring to address the worst of them, and a meaningful expansion of test coverage so that the next round of feature work doesn't reintroduce the same problems.

The founder becomes the bottleneck

The challenge: At MVP, the founder being in every loop was an asset.

At Launch, as support volume grows, product decisions stack up, and operational complexity multiplies, that same instinct becomes the constraint.

The transition from doing the work to designing the systems that do the work is one of the hardest shifts in the startup lifecycle.

Because there's rarely a clear moment when it happens, the risk is to miss it entirely and stay in builder mode while the organization stalls around you.

Telltale signs that this is happening include decisions that should take an hour now take a week for you to get around to them, support requests that pile up because only you know the answer, and operational tasks that only happen when you personally remember to do them.

The remedy is an all-out audit of everything you're personally handling, from the tiniest task to the most high-stakes decisions, in order to identify what can be systematized, what can be delegated, and what genuinely still merits founder time and attention.

The 创始人 becomes the bottleneck

The challenge: At 最小可行产品 (MVP), the 创始人 being in every loop was an asset.

At Launch, as support volume grows, product decisions stack up, and operational complexity multiplies, that same instinct becomes the constraint.

The transition from doing the work to designing the systems that do the work is one of the hardest shifts in the 创业公司 lifecycle.

Because there's rarely a clear moment when it happens, the risk is to miss it entirely and stay in builder mode while the organization stalls around you.

Telltale signs that this is happening include decisions that should take an hour now take a week for you to get around to them, support requests that pile up because only you know the answer, and operational tasks that only happen when you personally remember to do them.

The remedy is an all-out audit of everything you're personally handling, from the tiniest task to the most high-stakes decisions, in order to identify what can be systematized, what can be delegated, and what genuinely still merits 创始人 time and attention.

Security and compliance are no longer deferrable

The challenge: Keeping security and compliance measures simple was OK for MVP but now, with real users, real data, and potentially enterprise contracts on the table, it becomes a liability.

At MVP, with a handful of beta users and no sensitive data in production, security vulnerabilities were theoretical risks.

The hypothetical, however, becomes very real exposure risk the moment your product enters production with real users depending on it.

Furthermore, compliance requirements that didn't apply to a prototype definitely apply the moment you're handling customer data, processing payments, or selling into regulated industries.

The remedy is a systematic security and compliance review before production scale arrives, not after, and treat everything that surfaces as a required remediation—not a suggestion—before the next wave of users arrives.

Security and compliance are no longer deferrable

The challenge: Keeping security and compliance measures simple was OK for 最小可行产品 (MVP) but now, with real users, real data, and potentially enterprise contracts on the table, it becomes a liability.

At 最小可行产品 (MVP), with a handful of beta users and no sensitive data in 生产环境, security vulnerabilities were theoretical risks.

The hypothetical, however, becomes very real exposure risk the moment your product enters 生产环境 with real users depending on it.

Furthermore, compliance requirements that didn't apply to a prototype definitely apply the moment you're handling customer data, processing payments, or selling into regulated industries.

The remedy is a systematic security and compliance review before 生产环境 规模化 arrives, not after, and treat everything that surfaces as a required remediation—not a suggestion—before the next wave of users arrives.

Expansion before you're ready

The challenge: New markets and funding opportunities look like growth opportunities.

They can also be where product-market fit goes to die.

The initial traction you've built is real, but it's also specific to your early audience.

Expanding too early into a market that's meaningfully different from your original one introduces new user behaviors, compliance requirements, payment infrastructure, and baseline expectations that your product wasn't designed around.

Suddenly there are too many new variables and you lose the ability to interpret your own data clearly.

You also run the risk of neglecting your original user base while chasing a new and unproven audience.

How Claude can help Launch stage founders All three forms of Claude are in full use in the Launch stage, and they support each other: each tool produces outputs that become inputs for the other two.

The results compound organically, and a founder using all three tools together gets more than the sum of their parts.

This is what makes the ultra-lean startup model structurally possible.

When Claude Code builds the product, Claude Cowork builds the company around it, and Claude helps operationalize this product and organizational knowledge, a small team can run like a company nx its size.

Expansion before you're ready

The challenge: New markets and funding opportunities look like growth opportunities.

They can also be where 产品-市场契合 goes to die.

The initial traction you've built is real, but it's also specific to your early audience.

Expanding too early into a market that's meaningfully different from your original one introduces new user behaviors, compliance requirements, payment 基础设施, and baseline expectations that your product wasn't designed around.

Suddenly there are too many new variables and you lose the ability to interpret your own data clearly.

You also run the risk of neglecting your original user base while chasing a new and unproven audience.

How Claude can help 发布阶段 创始人s All three forms of Claude are in full use in the 发布阶段, and they support each other: each tool produces outputs that become inputs for the other two.

The results compound organically, and a 创始人 using all three tools together gets more than the sum of their parts.

This is what makes the ultra-lean 创业公司 model structurally possible.

When Claude Code builds the product, Claude Cowork builds the company around it, and Claude helps operationalize this product and organizational knowledge, a small team can run like a company nx its size.

Remediate technical debt before it compounds

Your MVP codebase works, but it also needs a systematic remediation pass in search of any technical debt that could become a structural liability.

First, use Claude Code to run a full architectural audit: identify where the codebase is brittle, any shortcuts that will become expensive to maintain, and where test coverage is thin enough that the next round of feature work will reintroduce the same problems.

Feed Claude Code's audit findings back to Claude to triage and sequence the remediation work: what needs to be fixed before the next release, what can wait a sprint, and what represents acceptable ongoing debt given your current stage.

This is also the moment to document the architectural decisions you made during the MVP stage (the ones that lived in your head because there was no time to write them down).

Getting them into a CLAUDE.

md now ensures that every future Claude Code session starts from a shared understanding of how the system was designed and why.

- Exercise: Direct Claude Code to audit your MVP codebase and produce a

prioritized list of structural weaknesses, test coverage gaps, and refactoring candidates.

Then feed that list to Claude and ask it to sequence the remediation work across your several sprints: any significant issues that you need to address first, things that can be handled in parallel with feature development, and things that can wait.

Remediate technical debt before it compounds

Your 最小可行产品 (MVP) codebase works, but it also needs a systematic remediation pass in search of any technical debt that could become a structural liability.

First, use Claude Code to run a full architectural audit: identify where the codebase is brittle, any shortcuts that will become expensive to maintain, and where test coverage is thin enough that the next round of feature work will reintroduce the same problems.

Feed Claude Code's audit findings back to Claude to triage and sequence the remediation work: what needs to be fixed before the next release, what can wait a sprint, and what represents acceptable ongoing debt given your current stage.

This is also the moment to document the architectural decisions you made during the 最小可行产品 (MVP) stage (the ones that lived in your head because there was no time to write them down).

Getting them into a CLAUDE.

md now ensures that every future Claude Code session starts from a shared understanding of how the system was designed and why.

Exercise: Direct Claude Code to audit your 最小可行产品 (MVP) codebase and produce a prioritized list of structural weaknesses, test coverage gaps, and refactoring candidates.

Then feed that list to Claude and ask it to sequence the remediation work across your several sprints: any significant issues that you need to address first, things that can be handled in parallel with feature development, and things that can wait.

Build the systems that replace founder attention

Building the operational systems that free your attention to handle responsibilities only the founder can tackle requires knowing exactly where your attention is going.

Use Claude Cowork to run a structured audit of your current operational load, documenting every recurring task, every decision that lands on your desk, and every workflow that only happens because you personally remember to do it.

Then have Claude Cowork categorize this inventory into what can be automated entirely, what needs a human but not necessarily you, and what genuinely requires founder judgment.

Once the audit is complete, use Claude Cowork to design the workflow logic for the automation candidates: what triggers each workflow, what the decision rules are, what the output looks like, and where it goes when it's done.

Build the systems that replace 创始人 attention

Building the operational systems that free your attention to handle responsibilities only the 创始人 can tackle requires knowing exactly where your attention is going.

Use Claude Cowork to run a structured audit of your current operational load, documenting every recurring task, every decision that lands on your desk, and every workflow that only happens because you personally remember to do it.

Then have Claude Cowork categorize this inventory into what can be automated entirely, what needs a human but not necessarily you, and what genuinely requires 创始人 judgment.

Once the audit is complete, use Claude Cowork to design the workflow logic for the automation candidates: what triggers each workflow, what the decision rules are, what the output looks like, and where it goes when it's done.

Make security and compliance a product workstream

Use Claude Code to surface code-level issues that frequently come up in SOC 2, GDPR, or HIPAA audits and standards that your target market requires.

This will surface both vulnerabilities and compliance gaps.

Feed those findings to Claude to help you prioritize the remediation work and design the controls, audit logging, and access management that enterprise buyers will ask for before they sign.

Note: AI scans are an aid but not a substitute for qualified compliance review.

Next, build the compliance workstream into your development cycle rather than running it as a one-time project; compliance documentation needs to be continually maintained and updated.

For founders approaching enterprise contracts or international markets, this is also the moment where the Claude Code security scan can help you prepare for an independent security assessment.

- Exercise: Run a code-level security review with Claude Code oriented to the

frameworks your target market requires.

Feed the output to Claude and ask it to produce two things: a prioritized security remediation sequence and a list of the documentation and controls you'll need to produce to satisfy a compliance review from a prospective enterprise buyer.

Make security and compliance a product workstream

Use Claude Code to surface code-level issues that frequently come up in SOC 2, GDPR, or HIPAA audits and standards that your target market requires.

This will surface both vulnerabilities and compliance gaps.

Feed those findings to Claude to help you prioritize the remediation work and design the controls, audit logging, and access management that enterprise buyers will ask for before they sign.

Note: AI scans are an aid but not a substitute for qualified compliance review.

Next, build the compliance workstream into your development cycle rather than running it as a one-time project; compliance documentation needs to be continually maintained and updated.

For 创始人s approaching enterprise contracts or international markets, this is also the moment where the Claude Code security scan can help you prepare for an independent security assessment.

- Exercise: Run a code-level security review with Claude Code oriented to the

frameworks your target market requires.

Feed the output to Claude and ask it to produce two things: a prioritized security remediation sequence and a list of the documentation and controls you'll need to produce to satisfy a compliance review from a prospective enterprise buyer.

Stand up the product management processes

you've been skipping The Launch stage requires a set of lightweight, repeatable processes that can run without requiring founder intervention to trigger or function.

Use Claude to design how your product timeline and work cycles will be structured, what a spec needs to include before Claude Code touches a feature, how bug reports get triaged and routed, and what your weekly metrics report covers and how it's distributed.

Once process design is done, use Claude Cowork to build and run the operational layer: scheduling sprint ceremonies, routing incoming bug reports to the right place, compiling weekly metrics from your connected data sources, and maintaining the feedback loop that keeps user signals flowing into product decisions.

- Exercise: Ask Claude to design a lightweight product management

operating system: a defined sprint cadence, a minimum spec template, a bug triage decision tree, and a weekly metrics brief that pulls from your actual data sources.

Then set up Claude Cowork to implement and run the system's recurring operational elements, like scheduling, routing, and report compilation, to happen on schedule without you.

Stand up the product management processes

you've been skipping The 发布阶段 requires a set of lightweight, repeatable processes that can run without requiring 创始人 intervention to trigger or function.

Use Claude to design how your product timeline and work cycles will be structured, what a spec needs to include before Claude Code touches a feature, how bug reports get triaged and routed, and what your weekly 指标 report covers and how it's distributed.

Once process design is done, use Claude Cowork to build and run the operational layer: scheduling sprint ceremonies, routing incoming bug reports to the right place, compiling weekly 指标 from your connected data sources, and maintaining the 反馈循环 that keeps user signals flowing into product decisions.

- Exercise: Ask Claude to design a lightweight product management

operating system: a defined sprint cadence, a minimum spec template, a bug triage decision tree, and a weekly 指标 brief that pulls from your actual data sources.

Then set up Claude Cowork to implement and run the system's recurring operational elements, like scheduling, routing, and report compilation, to happen on schedule without you.

Scale stage

During the Scale phase, the founder's role re-centers from builder to public-facing executive.

The product is still central, but your personal day-to-day work becomes increasingly about the company itself.

Your attention must expand to new Scale-stage activities like analyst briefings and IPO roadshows even as you strive to maintain the lean, AI-centered structural advantage.

Scale stage goals The work of scaling technical infrastructure keeps on going, and is now joined by the work of scaling the organization itself and maturing it into a business.

At the scale stage you're looking at going from thousands of users to millions, and from one market to many.

At every prior stage, growth was something you could feel your way through by being close to users and adjusting course based on data from tight feedback loops plus a healthy dose of founder instinct.

Now, though, the goal is to build systematic growth that's sustained by mature organizational operations.

For an AI-native startup, your goal should be to build a defensible moat through accumulated depth, stemming from the expertise you've built into your product, your product's depth of integration with the other tools and platforms your users rely on, and the proprietary system data and workflows.

The founders who've been building consistently in one direction, on consistent infrastructure, now have something genuinely hard to replicate.

At this stage, public investors, analysts, regulators, enterprise procurement teams, and acquirers apply greater pressure—along with greater skepticism—because the stakes are higher now.

Your product and org have to withstand external scrutiny: not just the capabilities of what you've built, but the governance, compliance posture, financial controls, and strategic narrative that surround it.

规模化阶段

During the Scale phase, the 创始人's role re-centers from builder to public-facing executive.

The product is still central, but your personal day-to-day work becomes increasingly about the company itself.

Your attention must expand to new Scale-stage activities like analyst briefings and IPO roadshows even as you strive to maintain the lean, AI-centered structural advantage.

规模化阶段目标 The work of scaling technical 基础设施 keeps on going, and is now joined by the work of scaling the organization itself and maturing it into a business.

At the 规模化 stage you're looking at going from thousands of users to millions, and from one market to many.

At every prior stage, growth was something you could feel your way through by being close to users and adjusting course based on data from tight 反馈循环's plus a healthy dose of 创始人 instinct.

Now, though, the goal is to build systematic growth that's sustained by mature organizational operations.

For an AI原生 创业公司, your goal should be to build a defensible moat through accumulated depth, stemming from the expertise you've built into your product, your product's depth of 集成 with the other tools and platforms your users rely on, and the proprietary system data and workflows.

The 创始人s who've been building consistently in one direction, on consistent 基础设施, now have something genuinely hard to replicate.

At this stage, public investors, analysts, regulators, enterprise procurement teams, and acquirers apply greater pressure—along with greater skepticism—because the stakes are higher now.

Your product and org have to withstand external scrutiny: not just the capabilities of what you've built, but the governance, compliance posture, financial controls, and strategic narrative that surround it.

Scale stage exit criteria

The exit condition at Scale is no longer a single milestone but a threshold event: the company is sustainable even as the founder is, increasingly, not directly running day-to-day operations.

You've demonstrated systematic growth; built organizational governance and compliance infrastructure that satisfies the most demanding external reviewers; and have a solid answer to the question, "If a well-funded incumbent copied your product today, would your users stay?"

" In practice, this threshold will typically take one of three forms: sustainable profitability at a scale that no longer requires external capital, IPO-readiness, or acquisition.

All three require that your growth is systematic and auditable, your product moat stands up under scrutiny, and your organization is operationally mature and sustainable.

When this is true, congratulations are in order: your startup has gone from being a bet to being a business.

规模化阶段退出标准

The exit condition at Scale is no longer a single milestone but a threshold event: the company is sustainable even as the 创始人 is, increasingly, not directly running day-to-day operations.

You've demonstrated systematic growth; built organizational governance and compliance 基础设施 that satisfies the most demanding external reviewers; and have a solid answer to the question, "If a well-funded incumbent copied your product today, would your users stay?"

" In practice, this threshold will typically take one of three forms: sustainable 盈利 at a 规模化 that no longer requires external capital, IPO-readiness, or acquisition.

All three require that your growth is systematic and auditable, your product moat stands up under scrutiny, and your organization is operationally mature and sustainable.

When this is true, congratulations are in order: your 创业公司 has gone from being a bet to being a business.

Delegating the operational layer

The challenge: Scale-stage operational systems have to run reliably and sustainably without being babysat.

For a founder who has been hands-on since day one, that transition can be as much a psychological challenge as a structural one.

Your Launch stage work was creating the systems; in the Scale phase, it becomes (1) maturing these systems until they are fully trustworthy and (2) then actually trusting them.

This is harder than it sounds.

Even if you're a founder who delegates well it's not always obvious what to hand off and what to keep on your plate .

Hand off too much, too fast—especially to AI-automated systems—and critical decisions get made without crucial context that only the founder can provide.

Hold on too long, though, and you can become a bottleneck.

The fundamental challenge here is identifying the institutional knowledge that lives only in the founder's head or undocumented workflows, and then codifying it into systems that are documented, auditable and transferable.

Delegating the operational layer

The challenge: Scale-stage operational systems have to run reliably and sustainably without being babysat.

For a 创始人 who has been hands-on since day one, that transition can be as much a psychological challenge as a structural one.

Your 发布阶段 work was creating the systems; in the Scale phase, it becomes (1) maturing these systems until they are fully trustworthy and (2) then actually trusting them.

This is harder than it sounds.

Even if you're a 创始人 who delegates well it's not always obvious what to hand off and what to keep on your plate .

Hand off too much, too fast—especially to AI-automated systems—and critical decisions get made without crucial context that only the 创始人 can provide.

Hold on too long, though, and you can become a bottleneck.

The fundamental challenge here is identifying the institutional knowledge that lives only in the 创始人's head or undocumented workflows, and then codifying it into systems that are documented, auditable and transferable.

Scaling technical operations

The challenge: Customers no longer evaluate only your product; they want to know that your organization can be a dependable infrastructure partner.

Technical challenges during the first three startup stages centered on the codebase: building the right solution without accruing technical debt and then hardening security and compliance for real users.

Having reached the Scale phase, the challenge now becomes everything built around the codebase; creating the support infrastructure, documentation, and reliability guarantees that signal maturity.

Larger-scale customers and institutional buyers signing multi-year contracts want these before they'll sign, and they'll also hold you to them once they do.

The same AI infrastructure that got you this far, though, helps you build dedicated support functions with defined response times and documentation that a new customer's engineering team can actually use.

Scaling technical operations

The challenge: Customers no longer evaluate only your product; they want to know that your organization can be a dependable 基础设施 partner.

Technical challenges during the first three 创业公司 stages centered on the codebase: building the right solution without accruing technical debt and then hardening security and compliance for real users.

Having reached the Scale phase, the challenge now becomes everything built around the codebase; creating the support 基础设施, documentation, and reliability guarantees that signal maturity.

Larger- 规模化 customers and institutional buyers signing multi-year contracts want these before they'll sign, and they'll also hold you to them once they do.

The same AI 基础设施 that got you this far, though, helps you build dedicated support functions with defined response times and documentation that a new customer's engineering team can actually use.

Scaling organizational functions

The challenge: A Scale-stage company generally needs organizational infrastructure like hiring, payroll, accounting, and legal operations, regardless of how many people are running it.

At Launch, systematizing operations meant automating the workflows consuming founder attention.

A Scale-stage startup now needs to grow an even broader, and in some ways more consequential, array of operational functions such as financial reporting, compliance monitoring, contract management, and customer support, to name a few.

Scaling organizational functions

The challenge: A Scale-stage company generally needs organizational 基础设施 like hiring, payroll, accounting, and legal operations, regardless of how many people are running it.

At Launch, systematizing operations meant automating the workflows consuming 创始人 attention.

A Scale-stage 创业公司 now needs to grow an even broader, and in some ways more consequential, array of operational functions such as financial reporting, compliance monitoring, contract management, and customer support, to name a few.

Building a GTM function

The challenge: Organic growth has a ceiling, and most Scale-stage founders hit it before they've ever had to build a real go-to-market function.

Idea, MVP, and Launch stage growth often originates from founder-led selling, from a well-timed Product Hunt post to personal relationships with early customers.

Organic growth like this works only to a certain point, though, and most startups hit this limit in the Scale phase.

Signs include flattening user curves, rising customer acquisition costs, and a pipeline that only moves when the founder is personally involved.

Scale-stage growth requires building a dedicated growth engine to reach new and broader audiences for your product.

Most startup founders, though, probably have never had to run things like marketing, sales, and analyst relations programs before.

A legit GTM motion requires not just establishing new systems and processes, but also creating a brand voice and story for how you want to talk about your product.

Because, at this stage in the startup lifecycle, you're going to need one to reach not only individual new users, but also entire target audiences like investors and enterprise buyers.

Fortunately, the GTM function doesn't have to be large to be effective, and the same AI infrastructure that built the product can bootstrap bringing it to market.

How Claude can help Scale stage founders Early startup stages use Claude as foundational infrastructure for the product itself: a research partner for validating the idea, the engineering team that designs and builds the prototype, and the AI operational layer that makes a single-founder startup possible.

AI-native startup founders who reach the Scale stage can now use Claude, Claude Code, and Claude Cowork to keep scaling the same way they built.

Building a GTM function

The challenge: Organic growth has a ceiling, and most Scale-stage 创始人s hit it before they've ever had to build a real go-to-market function.

Idea, 最小可行产品 (MVP), and 发布阶段 growth often originates from 创始人-led selling, from a well-timed Product Hunt post to personal relationships with early customers.

Organic growth like this works only to a certain point, though, and most 创业公司s hit this limit in the Scale phase.

Signs include flattening user curves, rising customer acquisition costs, and a 流水线/管道 that only moves when the 创始人 is personally involved.

Scale-stage growth requires building a dedicated growth engine to reach new and broader audiences for your product.

Most 创业公司 创始人s, though, probably have never had to run things like marketing, sales, and analyst relations programs before.

A legit GTM motion requires not just establishing new systems and processes, but also creating a brand voice and story for how you want to talk about your product.

Because, at this stage in the 创业公司 lifecycle, you're going to need one to reach not only individual new users, but also entire target audiences like investors and enterprise buyers.

Fortunately, the GTM function doesn't have to be large to be effective, and the same AI 基础设施 that built the product can bootstrap bringing it to market.

How Claude can help 规模化阶段 创始人s Early 创业公司 stages use Claude as foundational 基础设施 for the product itself: a 研究/调研 partner for validating the idea, the engineering team that designs and builds the prototype, and the AI operational layer that makes a single-创始人 创业公司 possible.

AI原生 创业公司 创始人s who reach the 规模化阶段 can now use Claude, Claude Code, and Claude Cowork to keep scaling the same way they built.

Handing off day-to-day tasks to Claude Cowork

Start the Scale stage with a clear-eyed view of where you most need to invest your time and attention now, which can be a challenge for first time founders who've never built a business before.

Claude can help by building the list of things only you should be doing at this stage, which could include things like product narrative decisions, board relationships, enterprise deals, and founder- to-founder conversations. Anything not on that list is a candidate for delegation or Claude Cowork automation.

- Exercise: Use Claude to produce a bottleneck map of your current operational

layer: every workflow, decision, and approval currently routed through you.

Now, ask Claude to extrapolate what happens to each one when you're unavailable for a week.

The workflows that stall are the ones where you are still hands-on enough to derail progress.

How do these map to the inventory of founder priorities and responsibilities you made with Claude?

Next, it's time to pressure-test that the systems you've already built are actually ready to scale with your business as it grows.

- Exercise: Use Claude to map your current workflows, and then ask it what

happens to each one when you're unavailable for a week.

The workflows that stall are the ones where handoff criteria, escalation paths, or exception handling still need tightening.

Claude can help analyze the failure points and recommend appropriate fixes so you can update or replace Claude Cowork automations as necessary.

Handing off day-to-day tasks to Claude Cowork

Start the 规模化阶段 with a clear-eyed view of where you most need to invest your time and attention now, which can be a challenge for first time 创始人s who've never built a business before.

Claude can help by building the list of things only you should be doing at this stage, which could include things like product narrative decisions, board relationships, enterprise deals, and 创始人- to-创始人 conversations. Anything not on that list is a candidate for delegation or Claude Cowork automation.

- Exercise: Use Claude to produce a bottleneck map of your current operational

layer: every workflow, decision, and approval currently routed through you.

Now, ask Claude to extrapolate what happens to each one when you're unavailable for a week.

The workflows that stall are the ones where you are still hands-on enough to derail progress.

How do these map to the inventory of 创始人 priorities and responsibilities you made with Claude?

Next, it's time to pressure-test that the systems you've already built are actually ready to 规模化 with your business as it grows.

- Exercise: Use Claude to map your current workflows, and then ask it what

happens to each one when you're unavailable for a week.

The workflows that stall are the ones where handoff criteria, escalation paths, or exception handling still need tightening.

Claude can help analyze the failure points and recommend appropriate fixes so you can update or replace Claude Cowork automations as necessary.

Scale technical operations into enterprise-grade infrastructure

As you scale, buyers need reassurance that your product and your organization can be trusted as long-term infrastructure.

Technical work still goes on inside the codebase as always, but now there is technical work around the codebase to handle, too.

The first step is to convert institutional knowledge into a system that scales.

Use Claude to draft and maintain the written infrastructure that enterprise procurement expects to see, including product documentation, support playbooks, and SLAs.

In parallel, direct Claude Code to audit and harden the codebase against the specific reliability and security standards that enterprise contracts require, and to build out the technical support infrastructure that Discord-based community support never had to provide: logging, monitoring, incident response tooling, and the observability layer that makes SLAs actually enforceable.

Claude Cowork then runs the operational layer of enterprise support itself: ticket routing, escalation workflows, documentation updates triggered by product changes, renewal tracking, and the reporting cadences that enterprise customer success relies on.

Together, these three give a small team the support posture of a much larger organization, which is exactly what signing a multi-year enterprise contract requires you to demonstrate.

- Exercise: Pick your three most demanding prospects or identify three ideal

customers for your product that you'd love to sign.

Ask Claude to produce a gap analysis: what documentation, SLAs, and support infrastructure would an enterprise procurement team at each of these accounts expect to see before signing a multi-year contract, and where do you currently fall short?

Use the output to sequence the technical and documentation work across Claude Code and Claude Cowork.

Scale technical operations into enterprise-grade 基础设施

As you 规模化, buyers need reassurance that your product and your organization can be trusted as long-term 基础设施.

Technical work still goes on inside the codebase as always, but now there is technical work around the codebase to handle, too.

The first step is to convert institutional knowledge into a system that 规模化s.

Use Claude to draft and maintain the written 基础设施 that enterprise procurement expects to see, including product documentation, support playbooks, and SLAs.

In parallel, direct Claude Code to audit and harden the codebase against the specific reliability and security standards that enterprise contracts require, and to build out the technical support 基础设施 that Discord-based community support never had to provide: logging, monitoring, incident response tooling, and the observability layer that makes SLAs actually enforceable.

Claude Cowork then runs the operational layer of enterprise support itself: ticket routing, escalation workflows, documentation updates triggered by product changes, renewal tracking, and the reporting cadences that enterprise customer success relies on.

Together, these three give a small team the support posture of a much larger organization, which is exactly what signing a multi-year enterprise contract requires you to demonstrate.

- Exercise: Pick your three most demanding prospects or identify three ideal

customers for your product that you'd love to sign.

Ask Claude to produce a gap analysis: what documentation, SLAs, and support 基础设施 would an enterprise procurement team at each of these accounts expect to see before signing a multi-year contract, and where do you currently fall short?

Use the output to sequence the technical and documentation work across Claude Code and Claude Cowork.

Build a real GTM function

Founder hustle got you this far, but scaling your startup requires creating and implementing an actual go-to-market strategy.

AI can help you build, then and run, that complete GTM engine.

Claude can assist with building foundational GTM resources from scratch: market segmentation, messaging architecture, analyst relations strategy, sales playbooks, and the investor-facing metrics narratives that matter once you're talking to public investors, enterprise buyers, and Wall Street analysts.

Each of these audiences has its own vocabulary and evaluates you against its own standards; Claude's job is to translate your product's value props into a product marketing approach that's relevant for each audience segment.

Now, Claude Cowork can become your tactical execution layer: content pipelines, outbound sequences, analyst briefing logistics, newsroom and PR cadences, CRM hygiene, pipeline reporting, and the many recurring cycles that turn GTM strategy into actual commercial motion.

Where the GTM motion requires product marketing infrastructure—interactive demo environments, integration documentation, sandbox tenants, API references, technical one-pagers—Claude Code can build it for you.

Buyers expect to evaluate your product technically and, in the Scale phase, a Loom video and a sales deck no longer suffice.

This is also the infrastructure that lets your GTM motion run asynchronously: a well-built demo environment closes deals while you're in board meetings.

Build a real GTM function

Founder hustle got you this far, but scaling your 创业公司 requires creating and implementing an actual go-to-market strategy.

AI can help you build, then and run, that complete GTM engine.

Claude can assist with building foundational GTM resources from scratch: market segmentation, messaging architecture, analyst relations strategy, sales playbooks, and the investor-facing 指标 narratives that matter once you're talking to public investors, enterprise buyers, and Wall Street analysts.

Each of these audiences has its own vocabulary and evaluates you against its own standards; Claude's job is to translate your product's value props into a product marketing approach that's relevant for each audience segment.

Now, Claude Cowork can become your tactical 执行 layer: content 流水线/管道s, outbound sequences, analyst briefing logistics, newsroom and PR cadences, CRM hygiene, 流水线/管道 reporting, and the many recurring cycles that turn GTM strategy into actual commercial motion.

Where the GTM motion requires product marketing 基础设施—interactive demo environments, 集成 documentation, sandbox tenants, API references, technical one-pagers—Claude Code can build it for you.

Buyers expect to evaluate your product technically and, in the Scale phase, a Loom video and a sales deck no longer suffice.

This is also the 基础设施 that lets your GTM motion run asynchronously: a well-built demo environment closes deals while you're in board meetings.

Turning domain expertise and institutional knowledge into AI context

Many ultra-lean startup founders are building highly specific apps or tools for a real-world problem they experience or observe first-hand in a particular sector.

Agentic AI now makes it possible for founders who have never written a line of code to use their domain expertise to build products that solve sophisticated problems.

Claude, Claude Code, and Claude Cework each play a part in converting founder knowledge into compounding product specificity.

Using Claude to capture, organize, and refine founder knowledge puts domain expertise somewhere the product can reach.

Through extended conversations, projects, and memory, a founder can share everything they know—industry jargon, regulatory gotchas, edge cases, frustrations, reasons why the obvious answers to this problem don't work—into a structured, searchable context.

Skills can then codify recurring workflows (e.g.

, "how I audit a commercial lease," "how I triage a patient intake form") into reusable routines Claude runs the same way every time.

Over months, this becomes a proprietary knowledge substrate that no generalist AI can match.

Externalizing your domain knowledge with Claude becomes invaluable for encoding industry-specific edge cases into your product: a generalist AI medical billing tool breaks on 340B drug program claims, for example, but yours has specific logic for them.

Claude Code helps you translate common frustrations experienced by other professionals in your field into validation logic, prompt refinements, or an MCP integration with a niche industry system your competitors haven't heard of.

As a result, your app or tool's depth and breadth both continually compound in a way that competitors simply can't replicate.

- Exercise: Identify one edge case a generic competitor would definitely get

wrong in your vertical.

Work with Claude Code to build a dedicated test case for it (not a unit test) based on a scenario you've actually seen.

Every time a similar edge case surfaces, add it. Your test suite becomes a map of your moat.

Turning domain expertise and institutional knowledge into AI context

Many ultra-lean 创业公司 创始人s are building highly specific apps or tools for a real-world problem they experience or observe first-hand in a particular sector.

Agentic AI now makes it possible for 创始人s who have never written a line of code to use their domain expertise to build products that solve sophisticated problems.

Claude, Claude Code, and Claude Cework each play a part in converting 创始人 knowledge into compounding product specificity.

Using Claude to capture, organize, and refine 创始人 knowledge puts domain expertise somewhere the product can reach.

Through extended conversations, projects, and memory, a 创始人 can share everything they know—industry jargon, regulatory gotchas, 边界情况s, frustrations, reasons why the obvious answers to this problem don't work—into a structured, searchable context.

Skills can then codify recurring workflows (e.g.

, "how I audit a commercial lease," "how I triage a patient intake form") into reusable routines Claude runs the same way every time.

Over months, this becomes a proprietary knowledge substrate that no generalist AI can match.

Externalizing your domain knowledge with Claude becomes invaluable for encoding industry-specific 边界情况s into your product: a generalist AI medical billing tool breaks on 340B drug program claims, for example, but yours has specific logic for them.

Claude Code helps you translate common frustrations experienced by other professionals in your field into 验证 logic, prompt refinements, or an MCP 集成 with a niche industry system your competitors haven't heard of.

As a result, your app or tool's depth and breadth both continually compound in a way that competitors simply can't replicate.

- Exercise: Identify one 边界情况 a generic competitor would definitely get

wrong in your vertical.

Work with Claude Code to build a dedicated test case for it (not a unit test) based on a scenario you've actually seen.

Every time a similar 边界情况 surfaces, add it.

Your test suite becomes a map of your moat.

Compound accumulated user data into a defensible advantage

As users interact with your product, they generate behavioral signals (i.e.

, which outputs they accept and which they reject), which informs the product roadmap.

Over time, you'll learn the specific patterns, preferences, and edge cases of your particular user base.

This is what we mean by compounding value: each improvement makes the product more useful, which drives more usage, which creates more feedback, which drives more improvement.

This data is time-locked, context-specific, and impossible for a copycat to recreate: you simply can't buy the behavioral fingerprint of thousands of users who've been refining their workflows inside your product.

Claude can help audit whatever user interaction data you've collected, identify the highest-signal behavioral patterns within it, and design the feedback loop that turns ongoing usage into systematic model improvement.

- Exercise: Feed Claude a summary of your product's interaction data: what

you've been collecting, how long you've been collecting it, and what you know about how users engage with your product over time.

Ask it to identify the three highest-signal behavioral patterns in that data and design a feedback loop that turns each one into a systematic model improvement.

Then ask it to help you draft a one-page moat narrative to inform product marketing: the story of how your data flywheel works, how long it's been spinning, and why a well- resourced competitor starting today couldn't replicate it in under two years.

Compound accumulated user data into a defensible advantage

As users interact with your product, they generate behavioral signals (i.

e.

, which outputs they accept and which they reject), which informs the product 路线图.

Over time, you'll learn the specific patterns, preferences, and 边界情况s of your particular user base.

This is what we mean by compounding value: each improvement makes the product more useful, which drives more usage, which creates more feedback, which drives more improvement.

This data is time-locked, context-specific, and impossible for a copycat to recreate: you simply can't buy the behavioral fingerprint of thousands of users who've been refining their workflows inside your product.

Claude can help audit whatever user interaction data you've collected, identify the highest-signal behavioral patterns within it, and design the 反馈循环 that turns ongoing usage into systematic model improvement.

- Exercise: Feed Claude a summary of your product's interaction data: what

you've been collecting, how long you've been collecting it, and what you know about how users engage with your product over time.

Ask it to identify the three highest-signal behavioral patterns in that data and design a 反馈循环 that turns each one into a systematic model improvement.

Then ask it to help you draft a one-page moat narrative to inform product marketing: the story of how your data flywheel works, how long it's been spinning, and why a well- resourced competitor starting today couldn't replicate it in under two years.

Create workflow lock-in

Compounding data network effects make your product harder to replicate, but user workflow lock-in makes your product harder to leave.

The longer users run your product inside their daily operations, the more deeply it gets embedded in how they actually work.

They've built automations on top of it, trained people to use it, and connected it to their data sources and other tools.

The prompts they've developed, the workflows they've refined, and the outputs they've standardized have all been shaped around what your product does and how it does it.

At this point, switching goes from product decision to full scale operational project.

The first step in creating workflow lock-in is asking Claude to map your current customer base by integration depth.

For each customer segment, identify what workflows they've built on top of your product and which integrations they depend on.

This shows where your product is sticking, and where it needs to go deeper.

The more integrations you offer, the more surface area a customer has to construct workflows that rely on your product.

Claude Code helps you quickly spin up native integrations with the data pipelines, project management tools, and other systems that your target users depend on.

- Exercise: Ask Claude to help you build a workflow integration audit for your

top ten customers.

For each one, document the automations they've built, the integrations they depend on, the team workflows that run through your product, and your estimate of their switching cost.

Then ask Claude to identify the patterns across the group: what types of integration create the deepest lock-in for your specific product, and what you could build or enable to deepen integration for customers who are currently at the surface.

Create workflow lock-in

Compounding data network effects make your product harder to replicate, but user workflow lock-in makes your product harder to leave.

The longer users run your product inside their daily operations, the more deeply it gets embedded in how they actually work.

They've built automations on top of it, trained people to use it, and connected it to their data sources and other tools.

The prompts they've developed, the workflows they've refined, and the outputs they've standardized have all been shaped around what your product does and how it does it.

At this point, switching goes from product decision to full 规模化 operational project.

The first step in creating workflow lock-in is asking Claude to map your current customer base by 集成 depth.

For each customer segment, identify what workflows they've built on top of your product and which 集成s they depend on.

This shows where your product is sticking, and where it needs to go deeper.

The more 集成s you offer, the more surface area a customer has to construct workflows that rely on your product.

Claude Code helps you quickly spin up native 集成s with the 数据管道s, project management tools, and other systems that your target users depend on.

- Exercise: Ask Claude to help you build a workflow 集成 audit for your

top ten customers.

For each one, document the automations they've built, the 集成s they depend on, the team workflows that run through your product, and your estimate of their switching cost.

Then ask Claude to identify the patterns across the group: what types of 集成 create the deepest lock-in for your specific product, and what you could build or enable to deepen 集成 for customers who are currently at the surface.

Same job, new rules

同样的工作，新的规则

Same job, new rules

In the AI area, the founder's job hasn't changed: find a real problem, build something that solves it, and scale it into a company that matters.

What's changed is the path to get there.

Across the four stages—Idea, MVP, Launch, and Scale—AI compresses quarters into weeks.

Validation cycles that used to take months now take afternoons.

A working prototype no longer requires a co-founder with the right stack; it requires a clear problem and a few focused sessions with a coding agent.

Launch readiness compresses from a pre-launch scramble into a continuous workstream.

And at scale, the operational weight that used to force early hires into firefighting roles can increasingly be handed off to AI, freeing your team to spend their attention on the judgment calls that become your moat.

The bottlenecks are no longer what you can build, but what you choose to build.

- Building AI Agents for Startups: Shares how startups use agents to become

less founder-dependent as they scale.

- Claude Code docs: Carries builders from initial installation to advanced agentic

workflows. Pro-tip: get started with the "How Claude Code works" overview.

- Claude Code best practices: Covers patterns that have worked inside

Anthropic and across engineering teams — context management, permissions, planning, and verification workflows.

- Using CLAUDE.md files: Walks through how to configure Claude Code for

your specific codebase. Essential reading for MVP-stage founders setting up their development environment.

- Claude Code power user tips: Highlights workflow patterns from the Claude

Code team itself, including parallel sessions and verification loops.

- Get started with Claude Cowork: Shares how teams can set up Claude

Cowork and start implementing skills, plugins, and other features that scale its impact across your startup.

- Tutorials: claude.com/resources/tutorials offers a searchable list of hands-on

walkthroughs for specific tasks. Founder stories

- How three YC startups built their companies with Claude Code: Examining

同样的工作，新的规则

In the AI area, the 创始人's job hasn't changed: find a real problem, build something that solves it, and 规模化 it into a company that matters.

What's changed is the path to get there.

Across the four stages—Idea, 最小可行产品 (MVP), Launch, and Scale—AI compresses quarters into weeks.

Validation cycles that used to take months now take afternoons.

A working prototype no longer requires a co-创始人 with the right stack; it requires a clear problem and a few focused sessions with a coding agent.

Launch readiness compresses from a pre-发布/上线 scramble into a continuous workstream.

And at 规模化, the operational weight that used to force early hires into firefighting roles can increasingly be handed off to AI, freeing your team to spend their attention on the judgment calls that become your moat.

The bottlenecks are no longer what you can build, but what you choose to build.

- Building AI Agents for Startups: Shares how 创业公司s use agents to become

less 创始人-dependent as they 规模化.

- Claude Code docs: Carries builders from initial installation to advanced agentic

workflows.

Pro-tip: get started with the "How Claude Code works" overview.

- Claude Code best practices: Covers patterns that have worked inside

Anthropic and across engineering teams — context management, permissions, planning, and verification workflows.

- Using CLAUDE.

your specific codebase.

Essential reading for 最小可行产品 (MVP)-stage 创始人s setting up their development environment.

- Claude Code power user tips: Highlights workflow patterns from the Claude

Code team itself, including parallel sessions and verification loops.

- Get started with Claude Cowork: Shares how teams can set up Claude

Cowork and start implementing skills, plugins, and other features that 规模化 its impact across your 创业公司.

- Tutorials: claude.

walkthroughs for specific tasks. Founder stories

- How three YC 创业公司s built their companies with Claude Code: Examining

how HumanLayer (F24), Ambral (W25), and Vulcan Technologies (S25) used Claude to get prototypes to market fast and scale AI-powered platforms with agentic coding workflows.

- GC AI's founders used domain expertise to build a responsive, Claude-powered

legal platform for how in-house teams actually work: company-specific playbooks, cross-functional stakeholders, and variable risk tolerance thresholds

- Carta Healthcare uses Claude to power their clinical abstraction platform,

processing 22,000 surgical cases per year and reducing data abstraction time by 66%.

- Anything, powered by Claude and the Agent SDK, has helped 1.5 million users

turn ideas into working software products without writing code, including a non-technical founder who built and is already selling a full recruiting platform.

Anything's AI agent handles the full build so solopreneurs can double down on their domain expertise.

- Cogent is an applied AI lab building agents to automate critical enterprise

security tasks.

The startup uses Claude as the reasoning layer for agents that automate investigation, prioritization, and remediation across the full vulnerability lifecycle.

- Airtree uses Claude Cowork as the center of its operations infrastructure,

uniting data that used to be scattered across a dozen different tools and teams.

Now, when one person builds a workflow automation with skills, everyone in the organization can use it to do all the things on their to-do list that never got done.

- Duvo builds AI agents that run procurement, supply chain, and category

management processes across ERPs, supplier portals, spreadsheets, email, and even phone calls.

Duvo is built entirely on Claude, using the Agent SDK to orchestrate across workflows.

- Zingage is an AI agent platform built for 24/7 automated operations for home-

care agencies.

The startup uses Claude's structured tool calling to orchestrate across an EMR and multiple communication channels, and Claude's contextual reasoning to build agents that can give nuanced, patient-tailored outcomes rather than pattern-matching to the most common response.

- Kindora is an AI-powered platform built by a nonprofit executive who used

Claude Sonnet to build a desperately-needed tool for intelligently matching charities with funders.

After filtering thousands of matches down to the few worth pursuing, Kindora's MCP connector lets nonprofits access its prospecting tools directly within Claude.

- Wordsmith was founded by a lawyer-turned-CTO to provide reliable AI-

powered legal technology for in-house legal teams.

how HumanLayer (F24), Ambral (W25), and Vulcan Technologies (S25) used Claude to get prototypes to market fast and 规模化 AI-powered platforms with 代理编程 workflows.

- GC AI's 创始人s used domain expertise to build a responsive, Claude-powered

legal platform for how in-house teams actually work: company-specific playbooks, cross-functional 利益相关者s, and variable risk tolerance thresholds

- Carta Healthcare uses Claude to power their clinical abstraction platform,

processing 22,000 surgical cases per year and reducing data abstraction time by 66%.

- Anything, powered by Claude and the Agent SDK, has helped 1.

turn ideas into working software products without writing code, including a non-technical 创始人 who built and is already selling a full recruiting platform.

Anything's AI agent handles the full build so solopreneurs can double down on their domain expertise.

- Cogent is an applied AI lab building agents to automate critical enterprise

security tasks.

The 创业公司 uses Claude as the reasoning layer for agents that automate investigation, prioritization, and remediation across the full vulnerability lifecycle.

- Airtree uses Claude Cowork as the center of its operations 基础设施,

uniting data that used to be scattered across a dozen different tools and teams.

Now, when one person builds a workflow自动化 with skills, everyone in the organization can use it to do all the things on their to-do list that never got done.

- Duvo builds AI agents that run procurement, supply chain, and category

management processes across ERPs, supplier portals, spreadsheets, email, and even phone calls.

Duvo is built entirely on Claude, using the Agent SDK to orchestrate across workflows.

- Zingage is an AI agent platform built for 24/7 automated operations for home-

care agencies.

The 创业公司 uses Claude's structured tool calling to orchestrate across an EMR and multiple communication channels, and Claude's contextual reasoning to build agents that can give nuanced, patient-tailored outcomes rather than pattern-matching to the most common response.

- Kindora is an AI-powered platform built by a nonprofit executive who used

Claude Sonnet to build a desperately-needed tool for intelligently matching charities with funders.

After filtering thousands of matches down to the few worth pursuing, Kindora's MCP connector lets nonprofits access its prospecting tools directly within Claude.

- Wordsmith was founded by a lawyer-turned-CTO to provide reliable AI-

powered legal technology for in-house legal teams.

Claude is the reasoning engine for Wordsmith's contract review, agreement drafting, and document review capabilities, and the startup's engineering team uses Claude Code for building and evolving the platform itself.

- Anthropic Startups Program: For startups working with Anthropic's VC

partners, the program provides free API credits, the highest tier of publicly available rate limits, and invitations to exclusive founder events and workshops.

- Claude community: Forums and community spaces for builders.
- Live learning resources: Conferences, webinars, livestreams, and recordings.

claude.ai

Claude is the reasoning engine for Wordsmith's contract review, agreement drafting, and document review capabilities, and the 创业公司's engineering team uses Claude Code for building and evolving the platform itself.

- Anthropic Startups Program: For 创业公司s working with Anthropic's VC

partners, the program provides free API credits, the highest tier of publicly available rate limits, and invitations to exclusive 创始人 events and workshops.

- Claude community: Forums and community spaces for builders.

- Live learning resources: Conferences, webinars, livestreams, and recordings.

claude.ai